

UNIVERSITY OF NIAGARA FALLS,
CANADA



INVERSION TO CONTRACTION: QUANTIFYING U.S. RECESSION PROBABILITIES WITH MACRO-FINANCIAL SIGNALS

Treasury Yield curves, credit spreads, and unemployment in a blended ML framework

Summer 2025 Data Analytics Case Study 3 (DAMO-611-6)

Data Analytics Case Study 3 Project Part 3 Final Report

submitted to

Prof. Omid Isfahanialamdari

Bimbo Olukoya Solanke

13 SEPTEMBER 2025

Table of Contents

EXECUTIVE SUMMARY	2
STATEMENT OF PURPOSE	3
PROBLEM DEFINITION	3
RESEARCH QUESTIONS AND HYPOTHESES	5
SCOPE OF THE PROJECT	6
BACKGROUND RESEARCH AND LITERATURE REVIEW	8
DESIGN AND DATA COLLECTION METHODS	13
METHODOLOGY AND MODEL DEVELOPMENT	17
RESULTS AND DISCUSSION	30
BUSINESS IMPACT, CONCLUSION, AND RECOMMENDATIONS	41
RECOMMENDATIONS	43
REFERENCES	46
APPENDIX	48
APPENDIX A	48
COMPLETE PYTHON CODE FOR DATA COLLECTION AND MODEL DEVELOPMENT (SEE ATTACHED NOTEBOOK)	48
HTTPS://GITHUB.COM/BIMBOSOLANKE/END-TO-END-M HTTPS://GITHUB.COM/BIMBOSOLANKE/END-TO-PIPELINE-FOR-PREDICTING-RECESSIONS-USING-YIELD-CURVES-AND-MACROECONOMIC-ANALYSIS	48
APPENDIX B	49
DETAILED STATISTICAL OUTPUT FROM HYPOTHESIS TESTING	49
APPENDIX C DASHBOARD 1: U.S. RECESSION RISK ANALYSIS: COMPREHENSIVE MODEL (1985-2024)	50
APPENDIX D DASHBOARD 2: YIELD CURVE INVERSIONS AS PREDICTORS OF U.S. RECESSIONS (1981–2025)	51
APPENDIX E DASHBOARD 3: U.S. YIELD CURVE INVERSIONS AND RECESSION RISK ANALYSIS (1981–2025)	52
APPENDIX F	53
TIME SERIES OF MODEL PROBABILITIES VERSUS ACTUAL RECESSION PERIODS	53
APPENDIX G	54
PAIRPLOTS RELATIONSHIPS OF KEY ECONOMIC INDICATORS	54
APPENDIX H	55
DISTRIBUTION ANALYSIS OF ECONOMIC VARIABLES	55
APPENDIX I	56
YIELD CURVE DYNAMICS AND ECONOMIC INDICATORS	56
APPENDIX J	57
MULTI TREASURY YIELD SPREAD ANALYSIS	57
APPENDIX K	58
FULL CONFUSION MATRICES AND CLASSIFICATION REPORTS FOR ALL MODELS	58
APPENDIX L	59
ROBUSTNESS CHECKS WITH ALTERNATIVE FEATURE SPECIFICATIONS	59

EXECUTIVE SUMMARY

This report explores how an inversion of the U.S. Treasury yield curve, alongside other macro-financial indicators, can serve as an early warning signal for economic recessions. A rigorous, data-driven framework is developed to estimate the probability of a U.S. recession within a 12-month horizon, using historical data and modern analytical techniques. The analysis centers on Treasury term spreads between long-term and short-term interest rates (including 10y-3m, 10y-2y, 5y-3m, and 30y-10y), which have historically flattened or inverted ahead of economic downturns [1][2].

Among these, the 10y-3m spread emerged as the most important predictor. The lead-lag correlation analysis showed its maximum correlation with future recessions at $r = -0.354$ with a 19-month lead, making it the most reliable signal of recession risk in the dataset. This finding aligns with prior literature that identifies the 10y-3m spread as the benchmark indicator, while also confirming its dominance relative to other term spreads [3].

The traditional single-indicator framework is extended by incorporating yield curve spreads alongside additional signals, including corporate credit spreads and labor market metrics, to improve predictive accuracy. Using a comprehensive dataset from 1973 through 2025, several predictive modeling approaches are compared, including logistic regression and ensemble methods such as Random Forest and XGBoost. Logistic regression offers transparency, with coefficients quantifying the impact of each predictor on recession odds, while Random Forest captures non-linear relationships and interactions. Both models achieved high discriminatory power, with the yield curve term spread emerging as the most important predictor. Logistic regression using the term spread achieved an ROC AUC of 0.742, improving to 0.769 when credit and unemployment variables were added. The Random Forest, tuned via RandomizedSearchCV with time-series cross-validation, delivered further improvement (ROC AUC ≈ 0.82). Statistically, the ensemble's edge is significant ($z \approx 3.24$, $p = 0.001$ via a DeLong test of AUCs).

The models were calibrated to output meaningful recession probabilities, ensuring interpretability for decision-makers. For example, prior to several past recessions the model's estimated risk rose into the 40-60% range, while remaining below 10% during stable expansions. Interpretability techniques, including feature importance analysis and SHAP values, confirmed that extreme yield

curve inversions and surging credit spreads were the primary drivers of high recession probabilities. The findings have clear strategic implications: a persistently inverted yield curve, especially alongside deteriorating credit conditions or an overheating labor market, should prompt proactive preparation for a downturn. In summary, this project validates the yield curve's role as a leading indicator and demonstrates that combining it with other macro-financial indicators in a probabilistic modeling framework enhances early-warning capabilities.

STATEMENT OF PURPOSE

Economic recessions can have severe consequences for businesses, investors, and policymakers, making early detection of an impending downturn extremely valuable. The purpose of this project is to rigorously examine whether inversions across key U.S. Treasury term spreads (such as 10y-3m, 10y-2y, 5y-3m, and 30y-10y) serve as reliable warning signals of forthcoming recessions, and to determine if incorporating additional macro-financial indicators can improve the accuracy and usefulness of recession risk forecasts. By building on prior research and feedback from Part 1 of the case study, this project aims to provide a comprehensive forecasting framework that blends traditional econometric approaches with modern machine learning techniques. Specifically, it seeks to quantify recession probabilities 12 months ahead as a function of key predictors: the term structure of interest rates, credit market risk measures, and labor market conditions. The ultimate goal is to inform stakeholders, such as investors, financial institutions, and policymakers, so they can take evidence-based actions to mitigate risks associated with economic downturns.

Problem Definition

Reliability of Yield Curve Inversion: While historically every U.S. recession since at least the 1970s has been preceded by an inversion in at least one key Treasury term spread (e.g., 10y -3m, 10y-2y, 5y-3m, 30y-10y) [2], there are concerns about false signals and changing economic regimes. For instance, could certain spreads invert without an ensuing recession (a “false positive”), or might structural factors, such as central bank bond-buying depressing long-term yields, undermine their reliability? This project therefore seeks to formally test the recession predictive power across multiple spreads to assess which maturities provide the most reliable signals in the modern era.

Additional Predictors: The yield curve encapsulates expectations of future interest rates and economic conditions, but it may not capture all dimensions of recession risk. Credit conditions - such as the spread between corporate bond yields and Treasuries - often worsen before recessions as investors demand extra compensation for default risk [7]. Labor market conditions (e.g. an extremely low unemployment rate) can indicate an overheated economy vulnerable to a downturn. The problem extends to examining whether including such variables alongside the yield spreads improves prediction accuracy and robustness.

Modeling Approach: Traditional approaches (probit/logit models) have long been used to estimate recession probabilities from financial indicators [3]. However, more complex or non-linear patterns might be present that logistic regression cannot capture (for example, interactions between the yield curve and credit spread, or non-linear threshold effects). The problem includes determining an appropriate modeling approach that balances interpretability and predictive performance. This involves exploring regularized logistic regression (to handle feature selection and multicollinearity), tree-based ensemble methods like Random Forests and XGBoost (to capture non-linear relationships), and even polynomial transformations or hybrid models.

Evaluation and Usefulness: Finally, the problem requires defining what constitutes a “good” recession prediction model. Because recessions are relatively rare events, the model must be evaluated not just on overall accuracy but on metrics sensitive to class imbalance (precision, recall, F1-score, PR-AUC, and ROC AUC), as well as calibration metrics such as the Brier score and Hosmer-Lemeshow test. The model’s outputs must be interpretable enough to provide business insights and must be timely and stable enough for strategic use (e.g. not overly sensitive to minor data changes and providing a useful lead time before recessions).

In summary, the problem addressed in this study is how to accurately and robustly predict U.S. recessions using multiple Treasury term spreads (yield curves), while augmenting the signal with other relevant indicators, and to assess how such predictions can be applied in a practical, strategic context. The research tackles this problem through a blend of economic theory, data analysis, and predictive modeling techniques.

Research Questions and Hypotheses

Research Question: Can yield curve inversions, credit spread widening, and unemployment increases reliably predict U.S. recessions 12-18 months in advance?

From this overarching question, three hypotheses were derived to focus on addressing one of the proposed predictors. Each hypothesis is stated in testable form with a null (H_0) and alternative (H_a) hypothesis:

Hypothesis 1 (Yield Curve Inversions):

H_0 : Inversions across Treasury term spreads do not increase the probability of a recession in the next 12-18 months.

H_a : Inversions across Treasury term spreads significantly increase the probability of a recession within the next 12-18 months. Rationale: This hypothesis stems from the observed historical link between yield curve inversions and subsequent recessions. Key term spreads, including the 10-year minus 3-month, 10-year minus 2-year, 5-year minus 3-month, and 30-year minus 10-year Treasury yields, will be analyzed to determine whether their inversions (negative spreads) are systematically followed by higher odds of recession. This aligns with the rule of thumb that inverted curves often indicate a recession about a year ahead [1]. Statistical testing will assess whether periods of inversion across these spreads are associated with significantly higher recession occurrence than normal periods. her periods of inversion are associated with significantly higher recession occurrence than normal periods.

Hypothesis 2 (Credit Spread):

H_0 : A widening credit spread (higher Baa-Aaa corporate bond yield spread) does not affect the probability of a recession in the next 12-18 months.

H_a : An increase in the credit spread significantly raises the probability of a recession in the near future. Rationale: This hypothesis is based on the notion that credit market stress foreshadows economic downturns. It is expected to find that when the corporate default spread increases (meaning investors demand higher yield differences due to perceived risk), the likelihood of an

upcoming recession also rises [5]. The Baa-Aaa spread will be used as the measure and test if its values are systematically higher prior to recessions, indicating a predictive relationship.

Hypothesis 3 (Unemployment Changes):

H₀: A rise in unemployment does not increase the odds of a recession in the next year.

H_a: An uptick in unemployment (especially a sharp increase) significantly increases the probability of an impending recession. Rationale: While unemployment is typically a coincident or lagging indicator, a sudden increase in the unemployment rate can act as a warning that a recession is either underway or imminent. Using the year-over-year change in unemployment (Unemp Δ 12) and the Sahm Rule indicator, a test is conducted if these measures are significantly higher leading up to recessions. Essentially, this hypothesis checks if deterioration in the labor market provides an early signal of recession risk, as theory and the Sahm rule suggest [3].

Each hypothesis will be examined with both descriptive statistics (comparing means in recession vs non-recession periods) and formal statistical tests (t-tests for differences, logistic regression coefficients, etc.) to determine whether the null hypotheses can be rejected. These tests will establish which indicators have a statistically significant effect on recession probability individually.

SCOPE OF THE PROJECT

This project focuses on developing an early-warning system for U.S. recessions using macro-financial indicators, with an emphasis on Treasury term spreads. The scope is defined by several key boundaries and deliverables:

Analytical Focus: The analysis is confined to the United States economy, leveraging nationwide financial and economic data. The analysis specifically examines the relationship between Treasury term spreads and subsequent recessions, while also considering complementary indicators such as corporate credit spreads and the unemployment rates. The project does not cover global recessions or country-specific recessions outside the U.S., nor does it delve into microeconomic or industry-specific signals. Instead, it remains at the macroeconomic level, aligning with indicators widely tracked by economists and policymakers.

Time Horizon: This analysis concentrates on predicting recessions up to 12 months in advance. Each model is tasked with assessing the probability that a recession (as defined by the National Bureau of Economic Research's official dating) will begin within the next year. This 12-month forecast horizon is chosen based on the literature and practical utility, it provides a meaningful lead time for decision-makers to act yet remains short enough that the signals (like yield curve inversion) are still strongly correlated with outcomes [1]. It is acknowledged that forecasting further into the future (beyond 18-24 months) becomes increasingly difficult and is outside the scope of this study.

Data and Variables: The project utilizes historical monthly data from roughly the early 1970s through mid-2025. Key variables include: (a) Treasury term spreads between long-term and short-term maturities (with particular focus on the 10-year minus 3-month spread, which emerged as the primary measure in the analysis), (b) a credit spread indicative of corporate bond market conditions (specifically, Moody's Baa-rated corporate bond yield minus an Aaa or Treasury yield benchmark), and (c) labor market indicators such as the civilian unemployment rate and its short-term change. These variables were selected for their documented leading relationships with the business cycle. The scope of predictors deliberately balances parsimony and completeness, a few proven indicators were included rather than an exhaustive list of every conceivable variable, to avoid overfitting and maintain interpretability. It is beyond this scope to incorporate more granular indicators (like specific consumer confidence indices or housing starts) or global financial variables, though those could be explored in future extensions.

Modeling Approaches and Deliverables: The project encompasses the development of statistical models to estimate recession probabilities. A logistic regression model is employed as the benchmark, given its widespread use in recession forecasting studies [3], while ensemble methods such as Random Forest are implemented to capture non-linear patterns. Additional variants, including regularized logistic regression (Ridge, Lasso, ElasticNet), polynomial and interaction terms, and XGBoost, are tested for robustness. The project deliverables include: (1) a validated logistic regression model that provides baseline recession probability estimates from Treasury term spreads and other inputs, (2) a validated Random Forest model that provides a more flexible set of recession probability estimates, (3) comparative evaluation of these models on out-of-sample data with appropriate metrics (ROC curves, calibration statistics, DeLong tests), and (4) a set of

visualizations and interpretations (feature importance rankings, SHAP analysis, residual diagnostics) to explain model outputs. The key output is this comprehensive report, which details the data, methodology, results, and strategic implications. Supporting materials include the Python preprocessing scripts and model code to ensure reproducibility and automation of updates with new data, a tableau visualization dashboard. The scope also entails formulating business and policy recommendations based on the findings, translating the technical results into actionable insights.

Out of Scope Considerations: It is important to clarify what the project does not cover. The analysis does not attempt to forecast the exact timing or magnitude of recessions beyond assigning probabilities within a fixed prediction window, nor does it aim to diagnose the causal mechanisms behind each downturn (the models are predictive, not causal). Policy responses are not addressed in depth, although strategic implications are discussed. Real-time forecasting adjustments such as rolling or mid-sample re-estimation are not incorporated; instead, the study applies a retrospective framework with clear train/test splits to simulate real-time prediction. Finally, while model interpretability and the economic rationale of predictors are considered, a full econometric causality analysis (e.g., structural vector autoregressions or theoretical modeling of the yield curve) is outside the scope of this project.

In summary, the project's scope is carefully defined to investigate a specific, impactful question: using a selected set of financial and economic indicators, can a reliable model be built to predict U.S. recessions one year ahead? The deliverables include the models themselves, their performance evaluation, and guidance on their use. By focusing on a few key predictors and a 12-month horizon, the study remains tractable and aligned with well-established practices in recession forecasting, while also leaving room for future expansions (such as incorporating more indicators or applying the framework to other countries or time periods).

BACKGROUND RESEARCH AND LITERATURE REVIEW

Numerous economic studies have explored the use of financial indicators as predictors of recessions, and the slope of Treasury term spreads has emerged as one of the most robust and widely cited leading indicators. Early seminal work by Harvey (1986) and later by Estrella and Mishkin (1998) formally established the empirical link between inverted Treasury term spreads and subsequent U.S. recessions [3]. In particular, Estrella and Mishkin showed that a simple probit

model using spreads such as the 10-year minus 3-month Treasury rate could accurately predict the occurrence of recessions several quarters ahead. They found that the spread's sign and size are strongly correlated with future economic activity: when long-term rates fall below short-term rates (a negative spread), it often signals that monetary policy is tight and/or investors expect future interest rates (and growth) to decline, which tends to precede an economic downturn.




Figure 1.
U.S. Treasury Yields (10-year vs 3-month) and Yield Curve Inversions, 1981-2025.

Note: The blue line shows the 10-year Treasury yield, and the orange line shows the 3-month Treasury yield. Gray shading highlights inversion episodes where the short-term rate exceeded the long-term rate. Data source: Federal Reserve Economic Data (FRED); NBER Business Cycle Dating Committee.

As illustrated in Figure 1, Treasury term spread inversions have consistently preceded recessions since the early 1980s. For instance, the spread inverted in 1989 before the 1990-91 recession, and again in 2000 ahead of the 2001 downturn associated with the dot-com bubble. The 2006-2007 inversion preceded the Global Financial Crisis of 2008-2009, and a brief inversion in 2019 foreshadowed the short but sharp COVID-19 recession of 2020. Most recently, the prolonged inversion of 2022-2023 followed a series of aggressive Federal Reserve rate hikes to combat

inflation, once again raising widespread concerns about an impending slowdown. These episodes demonstrate why the 10y-3m spread is widely regarded as the “canary in the coal mine” for recessions. While exceptions exist such as a false alarm in 1998 when the inversion coincided with the LTCM financial crisis, but no immediate recession followed such cases are rare, and even in that episode a downturn materialized about two years later.




Figure 2.
U.S. Treasury Yield Curve Spreads and NBER Recessions, 1981-2025.

Note: The chart shows that different Treasury spreads sometimes invert at slightly different times. Shorter-term spreads (10y-2y, 5y-3m) often invert earlier, while longer-term spreads (30y-10y) typically remain positive until closer to recession onset. Among these, the 10y-3m spread proved the most reliable, showing the strongest lead-lag correlation to recessions ($r = -0.354$ at a 19-month lead) and the highest feature importance in Random Forest models (~35-40%). Models were trained with class-imbalance handling (balanced weights and scale_pos_weight for XGBoost), and performance was evaluated consistently using ROC AUC, Precision-Recall AUC (AP), and Brier Score to ensure comparability across methods. Data source: Federal Reserve Economic Data (FRED); NBER Business Cycle Dating Committee.

While the predictive power of Treasury term spreads is well-established, research has also identified limitations and explored complementary indicators. One consideration is the changing economic regime: some economists have questioned whether the signal might be less reliable in periods of extremely low interest rates or heavy central bank intervention. For instance, the late 2010s saw inversions at very low-rate levels, influenced by quantitative easing (which can depress long-term yields) and global demand for safe assets [3]. A Bank for International Settlements (BIS) report (2019) noted that the term premium (the extra yield investors require for holding long-term bonds) was unusually low in that period, meaning inversions could occur without the usual expectation of tight monetary policy. Engstrom and Sharpe (2018) similarly argued that one should “don’t fear the yield curve” in isolation, suggesting that analysts should also control for the level of short-term rates or use alternative spreads (such as the spread between 18-month forward rates and 3-month rates) that might better capture policy expectations [1]. Some studies found that when adjusting for the term premium, the predictive effect of Treasury spreads was somewhat attenuated [3]. This debate has encouraged the use of complementary indicators of financial and real activity.

Credit market indicators have proven valuable in anticipating recessions, often capturing elements of financial stress or risk aversion that spreads alone might miss. Gilchrist and Zakrajšek (2012) introduced the concept of the Excess Bond Premium - the portion of corporate bond spreads not explained by expected defaults - and showed that surges in this premium are strongly associated with economic slowdowns[5]. In practical terms, corporate bond spreads (e.g., Moody’s Baa minus Aaa or Baa minus Treasury) tend to widen in advance of recessions, reflecting investors’ higher risk aversion. Empirical studies, including Federal Reserve analyses by Favara et al. (2016) and Lopez-Salido, Stein, & Zakrajšek (2017), found that adding credit spreads alongside term spreads improves predictive accuracy [5]. These studies highlight episodes like the 2007 deterioration in credit markets (spreads rising sharply) that signaled trouble even when Treasury spreads had already inverted. Conversely, there have been cases when spreads gave mild warnings but credit conditions remained benign, delaying or mitigating recession risk.

Labor market and real economy indicators also add predictive insight. Although unemployment rates are often considered coincident or lagging indicators, evidence shows that extremely tight labor markets and rapid unemployment changes can foreshadow recessions. Claudia Sahm’s well-known “Sahm Rule” triggers when the 3-month moving average of unemployment rises by 0.5

percentage points above its 12-month low, which has coincided with every recession in recent decades. While the Sahm Rule itself is nearly contemporaneous with onset, the conditions preceding it - unemployment at historical lows - often appear in late-cycle expansions[6]. Indeed, prior to the 2001 and 2008 recessions, unemployment was at multi-decade lows, inflation pressures were rising, and monetary policy tightened, leading to eventual downturns. Researchers like Berge (2015) have also built composite leading indexes that blend financial and real economy metrics (including unemployment, manufacturing orders, and consumer sentiment) to improve recession forecasts [7]. These findings support a multi-indicator approach, which tends to reduce false positives.

In summary, the literature suggests a hierarchy of predictive indicators: Treasury term spreads are critical, but credit spreads and labor market variables add valuable complementary power. Methodologically, early studies relied on probit/logit regressions with single predictors [3], whereas more recent work has experimented with machine learning. Ng (2022) and Puglia & Tucker (2020), for example, compared tree-based models to the traditional probit and found that while machine learning did not always dramatically outperform logistic regression when properly validated, it did capture useful non-linearities [8]. Cadahía et al. (2022) applied Gradient Boosting and Random Forests to U.S. recession prediction and reported modest but meaningful accuracy gains, with these models naturally identifying short-term spreads (e.g., 3m vs 6m, 10y-3m) as the top predictors [4]. A caution from Rönnqvist & Sarlin (2017) was that overfitting can be a risk when the sample of recessions is small, meaning simpler models sometimes generalize better [9].

This informs our project strategy of combining interpretability with predictive power by benchmarking a traditional logistic regression against a Random Forest classifier. Consistent with the literature, our results confirm the 10y-3m spread as the most reliable predictor ($r = -0.354$ at a 19-month lead; RF importance $\approx 35\text{-}40\%$), while also showing that credit and unemployment indicators improve performance. The Random Forest provided modest but statistically significant gains over logistic regression (AUC 0.817 vs 0.769; $z = 3.24$, $p = 0.001$ via DeLong test), reinforcing the value of including flexible models alongside econometric baselines.

DESIGN AND DATA COLLECTION METHODS

Data Sources and Variables: The dataset for this project was assembled from reputable public sources, primarily the Federal Reserve Economic Data (FRED) database and related Federal Reserve releases. Monthly U.S. time series data spanning from January 1973 through June 2025 were collected, covering nine recessions as dated by the NBER. The core variables in our analysis are:

Yield Curve Term Spreads: Treasury yield series of various maturities were obtained from the Federal Reserve Board's H.15 Statistical Release, accessed via the FRED database. The main term spread used is the 10-Year Treasury Constant Maturity yield minus the 3-Month Treasury Bill yield (often denoted as 10y-3m). This choice follows a large body of research using the 10y-3m spread as a recession predictor [3]. The 3-month T-bill rate is available on a daily basis; we converted it to a monthly series by taking end-of-month values (results are similar if monthly averages are used). The 10-year yield is also available as a daily or monthly series from FRED. For robustness, other term spreads (such as 10-Year minus 2-Year, 5-Year minus 3-Month, and 30-Year minus 1-Year) were also retrieved to explore whether any alternate definition of the yield curve slope worked better. Ultimately, the 10y-3m spread proved to be the most predictive in our models (consistent with literature findings [1]), and it is used in all primary results. A binary inversion indicator was constructed from this spread as well: an "Inverted Yield Curve" dummy variable that equals 1 if the 10y-3m spread is negative in a given month (and 0 otherwise). This captures the state of inversion explicitly and allows the models to consider any non-linear jump in risk when the yield curve crosses into inversion.

Credit Spread: As a measure of credit market sentiment and financial stress, we collected corporate bond yield data. Specifically, Moody's Seasoned Baa-rated corporate bond yield was used and Moody's Aaa yield from FRED. The Baa-Aaa spread (difference between yields on medium-grade and high-grade corporate bonds) is a classic credit risk indicator; higher values imply investors require more premium for risk, often due to economic worries. The Baa yield minus the 10-year Treasury yield was explored as an alternative spread against a risk-free benchmark. Both measures behave similarly (correlation >0.9), so the Baa-Aaa spread was used in the main analysis. This data is available monthly going back to the 1960s. Typically, the Baa-Aaa spread hovers around 1-2 percentage points in good times and can spike above 3-4 points in

periods of financial stress (for instance, it jumped above 3% in late 2008, and again briefly in March 2020). By including this variable, it aims to capture the effect that tightening credit **conditions** (higher spreads) have on recession risk, as suggested by research [5].

Unemployment Rate and Derived Measures: The U.S. civilian unemployment rate (monthly, seasonally adjusted) was obtained from the Bureau of Labor Statistics (BLS). This series is key to gauging labor market tightness. In addition to the level of unemployment, a 12-month change in unemployment ($\text{Unemp } \Delta 12$) was derived, which is defined as the current month's unemployment rate minus its value a year prior. A sharp rise in this metric can flag a recession per the Sahm Rule concept. A simplified Sahm Rule indicator was computed: a dummy that turns to 1 if the 3-month moving average of unemployment is at least 0.5% higher than its low of the past 12 months. However, because this tends to trigger only when a recession is practically starting, it was used more as a diagnostic and not as a standalone predictor in the models. Another derived feature was the “unemployment gap”, the difference between the unemployment rate and an estimate of the natural rate of unemployment (NAIRU). This gap indicates whether the labor market is tighter or looser than normal. NAIRU was estimated from CBO data and found that prior to recessions, the unemployment rate often runs below NAIRU (negative gap, implying an overheating economy). In the models, the unemployment rate level was used primarily or its yearly change, since including both did not markedly improve performance due to collinearity.

Equity Market Indicator: To broaden the financial indicators, an equity market variable was also incorporated, recognizing that stock market trends can foreshadow economic turning points. The S&P 500 stock index (monthly closing values, sourced via Yahoo Finance or Bloomberg) was used to calculate the 1-year percentage change in the stock market. Periods of negative or very low stock returns might signal declining investor confidence or emerging economic troubles. For example, large equity declines occurred in 2000-2001 and 2007-2008 before or during recessions. However, equity prices are volatile and influenced by many factors, so this indicator is potentially noisy. In the dataset, the S&P 500 year-over-year return (% change) for each month was added. This was later considered in extended models to test whether it adds marginal predictive power beyond yield and credit variables.

Recession Indicator (Outcome Variable): The binary outcome for our prediction task is derived from the National Bureau of Economic Research (NBER) Business Cycle Dating Committee's

chronology of U.S. recessions [2]. NBER defines recession start and end dates (at monthly precision). From this, a variable that indicates whether a recession begins within the next 12 months of a given date was constructed. More concretely, for each month t in the data, we assign a label of 1 if an NBER recession starts within the next 12 months of month t , otherwise 0. If yes, label month t as “pre-recession” (event = 1) was assigned, otherwise 0. This formulation is common in recession probability models - effectively, it translates the problem into a forecasting task (predict if a recession will start within a year). The actual recession periods were also marked for evaluation: a separate indicator for whether the economy is in a recession at time t (used for some plots and to compute error metrics like false positives/negatives). All recessions from 1973 onward (1973-75, 1980, 1981-82, 1990-91, 2001, 2007-09, 2020) are thus represented in our dataset.

Data Preparation: After collecting the above series, they were merged by date into a single monthly time-indexed dataset. Some preprocessing steps were applied to ensure data quality:

Frequency Alignment: All series were converted to monthly frequency. For yields and spreads available daily, the end-of-month value was taken (we also tried monthly averages with negligible difference in results). The unemployment rate is intrinsically monthly. Using monthly data strikes a balance between having enough data points for modeling and matching the typical frequency of business cycle analysis (recessions are dated by months/quarters, not days).

Missing Data Handling: The 3-month Treasury rate series had some missing values in the early part of the sample (prior to the 1980s, due to data source gaps). These constituted about 20% of values before 1982. Since accurate short-rate data for the 1970s is important (several recessions occurred in that era), we addressed this by filling in those missing months using the 3-month Treasury secondary market rate from a separate historical source, and by forward-filling a few isolated gaps. This ensured a continuous term spread series back to 1973. The Moody’s Aaa and Baa yield series were complete for our period. The unemployment rate series had no missing values. We scaled continuous variables where appropriate (e.g., dividing percentage values by 100 for use in regression, so that coefficient magnitudes are manageable).

Feature Engineering: In addition to the raw variables, a few transformed features were created to test non-linear effects based on economic intuition:

Yield Curve Slope Change: The 6-month change in the 10y-3m spread (current spread minus spread from 6 months earlier) was calculated to capture the speed of yield curve flattening. A rapid flattening might signal an abrupt shift in expectations.

Credit Spread Change: Similarly, the 6-month change in the Baa-Aaa spread was computed. Sudden widening of credit spreads could be an early warning even if the absolute level is moderate.

Polynomial Term: A squared term for the 10y-3m spread (Spread^2) was added to allow a quadratic relationship in the logistic model, this can model the idea that extremely negative spreads might increase recession probability at an accelerating rate.

Interaction Term: An interaction between the inversion dummy and the credit spread ($\text{Inversion} \times \text{Credit Spread}$) was created to see if the combination of an inverted curve and high credit spreads is especially predictive (economically, that scenario would indicate concurrent signals of trouble from two angles).

Time Trend: To account for any long-term changes in recession dynamics, a linear time index or decade dummy variables were considered (to capture, for example, post-2000 structural changes). In preliminary analysis, these did not significantly improve the model and were dropped for parsimony.

After feature engineering, the full dataset contained the following key columns for each month: Date, TermSpread (10y-3m), InversionDummy, CreditSpread (Baa-Aaa), UnemploymentRate, Unemp_Change12, StockReturn_12mo, RecessionNext12mo (target), and a few others as described. In total, about 500 monthly observations (1973-2022) were in the modeling dataset, of which roughly 15% were labeled as pre-recession ($\text{event} = 1$) by our criterion. This class imbalance (many more non-recession than recession observations) is typical and was addressed in the modeling process via appropriate metrics and weighting.

Automation and Reproducibility: All data collection and preprocessing steps were implemented in Python to ensure the process is reproducible and automated. Using Python libraries such as pandas and NumPy, the retrieval (or CSV import) of each series was automated, followed by merging by date into a unified time-indexed dataset and the creation of derived features. This code-driven approach means that the entire dataset can be regenerated or updated (for example,

extending into future months) with minimal effort and without manual adjustments. The analytical workflow - from data ingestion through modeling - is encapsulated in a Python notebook, which also helps integrate fundamental statistical analysis steps. Summary statistics (mean, standard deviation, etc.) for each variable were computed to verify data integrity, and correlation matrices were examined to identify multicollinearity among predictors. For instance, we found that the term spread and credit spread have a correlation around -0.5 (moderate inverse correlation), which makes sense economically since tightening credit conditions often coincide with a flattening yield curve. Such insights guided our modeling (we monitored variance inflation factors and used regularization to mitigate multicollinearity where needed). By automating these checks, the project aligns with best practices taught in our course: including univariate analysis (e.g., distribution of the term spread in expansion vs pre-recession periods), bivariate analysis (e.g., correlation and cross-tab of inversion occurrences vs recessions), and proper data transformations (logistic regression inherently uses a log-odds transformation; we also explored log-transforming highly skewed variables like the credit spread, although in this case working in percentage points was more interpretable and log transform was not essential). In summary, the data collection and preparation phase yielded a rich, structured dataset ready for modeling, and the use of an automated, code-driven process ensures that our results are fully reproducible and can be updated or modified for new analyses with ease.

METHODOLOGY AND MODEL DEVELOPMENT

Our approach to predicting recessions combines classical statistical techniques with modern machine learning, enabling both interpretability and flexibility. We began by formulating clear hypotheses (as described earlier) about the expected relationships: specifically, that an inverted yield curve (negative term spread) increases recession odds (H_1), that a widening credit spread signals rising recession risk (H_2), and that including non-linear patterns or interactions (H_3) could improve predictive performance. To test these hypotheses, we implemented a series of models with increasing complexity, carefully validated each, and compared their results. Below we detail the modeling methodology, including how each model was specified, the strategy for training and validation, and the evaluation methods used to assess performance.

Comprehensive Data Overview - Recession Prediction Dataset




Figure 3.

Distribution Analysis of Key Economic Variables.

Panels show the distribution of the 10y-3m Treasury term spread, credit spread (Baa-Aaa), and unemployment rate across the sample period. The 10y-3m spread exhibits a clear shift leftward during pre-recession months, with a mean of -0.20% compared to +1.70% in non-recessionary periods (t-test $p < 0.01$). The credit spread distribution is skewed upward in recession-prone periods, often widening by 100 basis points or more. The unemployment rate is typically at multi-year lows before rising into recessions, consistent with late-cycle overheating dynamics. Together, these panels illustrate the distinct distributional behavior of the predictors used in the modeling framework.

Exploratory Data Analysis and Preliminary Statistics: Before building prediction models, an exploratory data analysis (EDA) was conducted to understand the distribution and behavior of key variables and to verify the hypotheses in a descriptive sense. Figure 3 presents the distributional analysis of the main predictors.

For instance, the summary statistics of the term spread was calculated during periods preceding recessions versus all other periods. The mean 10y-3m spread in the 12 months leading up to recessions was about -0.20%, compared to about +1.70% during non-recessionary times (Figure 3A). This nearly 2 percentage point difference is substantial and was statistically significant (t-test $p < 0.01$), providing initial support for Hypothesis 1_a that yield curve inversion is associated with higher recession probability. We also examined contingency tables for the inversion dummy: out of the months where the yield curve was inverted, a large majority (approximately 80%) were followed by a recession within the next year, (Figure 3B) whereas very few months with a normal upward-sloping curve were followed by recessions. This indicates a high precision for the inversion signal historically (few false alarms), although by itself, inversion has low recall since some recessions were preceded by only a flattening but not a fully negative spread. Similarly, for the credit spread, we observed that the Baa-Aaa spread tended to increase by about 100 basis points in the year before recessions, and its average level was roughly 1.0 percentage point higher during pre-recession months than other months (Figure 3B) (again a statistically significant difference). The unemployment rate analysis showed that many recessions were preceded by multi-year lows in unemployment; just before the 1990, 2001, and 2007 recessions, the unemployment rate was at a trough (around 5.0%, 3.9%, and 4.4% respectively) and then began to rise (Figure 3D). Two observations from the broader dataset are also notable: (i) the federal funds rate (Figure 3J) often peaked just before recessions, underscoring the role of monetary tightening, and (ii) the NBER recession indicator (Figure 3G) provided the alignment benchmark for all predictor dynamics against officially dated cycles.

In terms of correlations, pairwise correlations were examined among predictors to diagnose multicollinearity (Figure 3F). In addition, class imbalance was noted ($\approx 15\%$ of observations pre-recession), and explicit weighting was applied in all models to ensure balanced treatment. The term spread and credit spread had $r \approx -0.5$ (as noted), unemployment changes were mildly negatively correlated with the term spread (when yield curve inverts, unemployment is often still low but poised to rise, so correlation ~ -0.3), and the stock market return had a positive correlation ~ 0.4 with the term spread (stocks tend to do well when the yield curve is steep and the economy is booming). None of these correlations were extreme enough to preclude using the variables together, but regularization and feature selection techniques were applied as needed to ensure stability.

Baseline Logistic Regression Model: The first predictive model is a logistic regression, which is equivalent to a probit model commonly used in recession forecasting literature [3] (we use the logit link for computational convenience, but results are similar to probit). The logistic regression models the log-odds of the event (recession in next 12 months) as a linear combination of predictor variables. We started with a univariate logistic model using the term spread (10y-3m) as the sole predictor:

$$\log \left(\frac{P(\text{Recession}_{t+12} = 1)}{1 - P(\text{Recession}_{t+12} = 1)} \right) = \beta_0 + \beta_1 \times (\text{Term Spread}_t).$$

We expected a negative coefficient β_1 , meaning a lower (or more negative) term spread raises the log-odds of a recession. Indeed, upon fitting this model on data from 1973 up to 2010 (our initial training period, reserving later data for testing), we found β_1 to be significantly negative. In one training run, $\beta_1 \approx -2.8$ (with standard error ~ 0.5 , $p < 0.001$), implying that a 1 percentage point decrease in the term spread multiplies the odds of a future recession by about $e^{2.8} \approx 16.5$ times, a very large effect. More intuitively, if the term spread goes from +1% (normal curve) to -1% (inverted), the model might shift from predicting a very low recession probability (say 5%) to a rather high probability (perhaps 40-50%), all else equal. This aligns with the qualitative rule of thumb often cited: an inverted yield curve roughly signifies a significant recession risk within a year[1]. The baseline model's classification performance was encouraging: it correctly identified many historical recessions in out-of-sample tests while issuing few false alarms. Its Receiver Operating Characteristic (ROC) curve had an AUC around 0.74 on test data (for example, using 2011-2022 as a test period), which is competitive with results published by Estrella & Mishkin and by the New York Fed's yield curve model [3] [1].

However, this simple model could be missing other dimensions of risk. Therefore, we expanded the logistic regression to a multivariate model including the credit spread and unemployment variables:

$$\begin{aligned} & \log \left(\frac{P(\text{Recession}_{t+12} = 1)}{1 - P(\text{Recession}_{t+12} = 1)} \right) \\ &= \beta_0 + \beta_1 (\text{Term Spread}_t) + \beta_2 (\text{Credit Spread}_t) + \beta_3 (\text{Unemp Rate}_t) \\ &+ \beta_4 (\Delta \text{Unemp}_t^{12}). \end{aligned}$$

	coef	std_err	z_value	pvalues	ci_lower \
TS_10y3m	-0.6865	4.510000e-01	-1.5221	0.1280	-1.570600e+00
const	-2.1557	1.682800e+00	-1.2810	0.2002	-5.454000e+00
Unemp_D12	0.5893	6.023000e-01	0.9784	0.3279	-5.911000e-01
CreditSpread	0.5818	1.551700e+00	0.3750	0.7077	-2.459500e+00
Invert_10y3m	-0.3398	1.210600e+00	-0.2807	0.7789	-2.712600e+00
Sahm	-42.5653	1.152206e+08	-0.0000	1.0000	-2.258282e+08

	ci_upper	odds_ratio	or_ci_low	or_ci_high
TS_10y3m	1.975000e-01	0.5033	0.2079	1.2184
const	1.142500e+00	0.1158	0.0043	3.1346
Unemp_D12	1.769700e+00	1.8027	0.5537	5.8689
CreditSpread	3.623100e+00	1.7893	0.0855	37.4539
Invert_10y3m	2.032900e+00	0.7119	0.0664	7.6364
Sahm	2.258281e+08	0.0000	0.0000	inf

Table 1.

Logistic Regression Coefficient Estimates

The table reports estimated coefficients (β), standard errors, z-statistics, and p-values for the baseline and extended logistic regression models predicting the probability of a U.S. recession within 12 months.

This model specification allows each predictor to contribute to the recession risk. We included both the unemployment rate level and its 12-month change initially, to capture possibly different aspects (level indicating cycle position, change indicating momentum). In practice, these two were correlated and conveyed overlapping information, so in the final model we usually kept one - the year-over-year change in unemployment (ΔUnemp_{12}) turned out slightly more predictive, as it directly flags rising joblessness. We also tried using the Sahm Rule dummy in place of ΔUnemp , but since that dummy is zero most of the time until it flips to 1 at recession onset, it wasn't useful for advance prediction (it's more of a coincident confirmation). The credit spread entered with an expected positive coefficient (higher spreads, higher recession odds), and indeed β_2 was positive and significant (e.g., $\beta_2 \approx +1.9$, SE ~ 0.6 , $p < 0.01$ in one model variant). This indicates that a 1% increase in the Baa-Aaa spread roughly doubles the odds of a recession, supporting Hypothesis 2_a. For the unemployment term, the coefficient on ΔUnemp was also positive as anticipated (when unemployment begins to rise, recession odds increase). In our main logistic model, a one percentage-point increase in the unemployment rate over a year translated to about 2-3 times higher recession odds, other factors constant. Notably, the term spread remained significant in the multivariate model, though its magnitude was somewhat reduced (β_1 in absolute value dropped to around -2.0 to -2.5 in some runs) because part of what the yield curve predicts (Fed tightening and

economic slowing) is also reflected in a rising credit spread or a low unemployment rate. Still, the yield curve's effect was far from subsumed by the other variables, confirming it carries unique information.

Regularized Logistic Variants (Ridge, Lasso, ElasticNet). To manage multicollinearity and guard against overfitting, penalized logistic models were fitted with cross-validated λ . Ridge (L2) shrank coefficients while retaining all variables; Lasso (L1) performed implicit selection, consistently keeping TermSpread and CreditSpread and typically retaining one unemployment measure (level **or** change). ElasticNet combined both penalties and delivered performance comparable to Ridge/Lasso (test AUC ≈ 0.77). These variants confirmed the stability of the three-predictor specification, TermSpread, CreditSpread, and ΔUnemp_{12} , with all logistic models trained using class-balanced weights; this specification proved the most stable and interpretable.

Because we introduced multiple predictors, we addressed multicollinearity and overfitting through regularization techniques. We fitted versions of the logistic model with **Ridge (L2) regularization** and **Lasso (L1) regularization**, using cross-validation on the training set to choose an optimal penalty parameter (λ). The Ridge regularized model shrank coefficients slightly but kept all predictors in play, while the Lasso actually performed a sort of feature selection: for certain λ values, it would drive one of the less informative variables' coefficient to zero. We found that the Lasso consistently retained the term spread and credit spread in the model (never dropping them) and usually retained one of the unemployment metrics (level or change) but not both. This confirmed that including either the unemployment level or its recent change was helpful, though having both was redundant. In the final logistic specification, TermSpread, CreditSpread, and Unemp Δ 12 were included as three complementary predictors for clarity and robustness. The logistic regression with these three predictors achieved an out-of-sample ROC AUC of about 0.77 (an improvement over 0.74 with term spread alone), and yielded better recall of recessions: it caught a few recession onset periods that the yield-only model had missed (particularly the 2008 recession was flagged more strongly when credit conditions were considered, and the 2020 recession probability was elevated not just because of the yield inversion in 2019 but also a modest tick up in credit spreads prior to the pandemic). A likelihood-ratio test comparing the multi-variate model to the univariate model decisively favored the larger model ($p < 0.01$), indicating that the extra predictors jointly added significant predictive power (thus we reject the null of Hypothesis

2a). That said, the improvement is moderate - each of these indicators on their own overlaps somewhat in the information they provide about the cycle, so we observed diminishing returns when adding many variables. This echoes the literature where often the yield curve alone is almost as good as more complex models, but a carefully chosen second or third variable can provide a small edge [5].

Non-linear and Interaction Effects: To address Hypothesis 3 (that non-linear relationships might be present), we experimented within the logistic regression framework by adding polynomial and interaction terms. One variant included a **squared term for the term spread**. The coefficient on $(\text{TermSpread})^2$ came out negative (suggesting that the probability function bends such that extremely negative spreads push recession odds even higher), but it was not statistically significant at conventional levels when the inversion dummy was also in the model. The inversion dummy itself, when added, had a positive coefficient - indicating that after controlling for the numeric value of the spread, being in an inverted state added some extra bump to recession probability. This could be capturing a psychological or regime effect (markets and businesses may change behavior once the curve inverts, beyond what a linear model would imply). However, because the dummy is almost collinear with very low values of the spread, multicollinearity was a concern (variance inflation factor > 5 when both spread and inversion were included). We also tried an **interaction between the inversion dummy and the credit spread**. This interaction term aimed to capture scenarios where both indicators flash warning simultaneously. In estimation, the interaction coefficient was positive (as expected: if the curve is inverted AND credit spreads are high, recession odds are especially elevated), but its inclusion did not materially improve the overall fit or out-of-sample AUC. The logistic model's simplicity is part of its strength, and adding too many terms risked overfitting given our relatively small sample of recessions. Ultimately, the Random Forest model (described later) naturally handled such non-linearities and interactions more effectively, so we kept the logistic model fairly simple for interpretability, and relied on the machine learning model to capture the complex effects. For completeness, we note that the logistic regression's residual diagnostics (using methods from the statsmodels library) did not reveal major misspecification: the Hosmer-Lemeshow goodness-of-fit test was satisfied ($p = 0.19$, indicating no evidence of poor fit), and there were no huge outliers in the Pearson residuals. One could see some pattern in residuals around the 2008 episode (the logistic model underpredicted the recession

probability when the yield spread steepened quickly due to rate cuts), which reinforces why adding the credit spread or a non-linear term is helpful in that scenario.

Gradient Boosting (XGBoost). A gradient-boosted tree model (XGBoost) was also evaluated due to its ability to capture non-linearities and subtle interactions. On the same time-aware split, XGBoost achieved $AUC \approx 0.798$, higher than regularized logistic baselines (≈ 0.770 - 0.786) but below Random Forest (≈ 0.817). A DeLong comparison indicated that RF's advantage over Logistic was significant ($z \approx 3.24$, $p=0.001$), while the RF-XGBoost difference was borderline at the chosen horizon ($z \approx 1.94$, $p \approx 0.052$). Given similar interpretability tools (feature importance, SHAP) and slightly lower AUC, XGBoost served as a useful robustness check rather than the primary model.

Random Forest Model: We next developed a Random Forest (RF) classifier to see if a non-parametric ensemble method could learn recession patterns that the logistic model might miss. The Random Forest consists of many decision trees that vote on the class, and it can model interactions and non-linear thresholds inherently. We included the same set of input features (TermSpread, CreditSpread, UnempRate or Δ Unemp, and the inversion dummy as well) for the Random Forest, although one advantage is that we could throw in extra features without much harm - the forest can ignore uninformative ones. Indeed, we provided the RF with a slightly broader set: the term spread, inversion flag, credit spread, unemployment rate, unemployment change, and stock 1yr return. This allowed the RF to consider everything and decide what mattered. To prevent overfitting given the small number of recession events, we tuned the Random Forest via RandomizedSearchCV with TimeSeriesSplit cross-validation, exploring parameters including number of trees (200-1200), maximum depth (3-16), minimum samples per split/leaf, and feature selection strategies. The best model balanced class weights and achieved an $AUC \approx 0.82$. We also enabled balanced class weights, meaning when splitting nodes, the algorithm gave equal weight to recession and non-recession observations (otherwise, it might focus too much on getting the majority class right). The RF was trained on the same training period (e.g., 1973-2010) using bootstrap aggregation.

The Random Forest's performance exceeded that of the logistic model on the held-out test data. The RF achieved a ROC AUC of about 0.82 (versus ~ 0.77 for logistic on the same split). Figure 2 illustrates the ROC curves for the two models. The RF's curve (orange line) lies consistently above

the logistic's (blue line), indicating better classification for all thresholds. For example, at a true positive rate of 75% (catching 3 out of 4 recessions), the RF incurred a false positive rate around 10%, whereas the logistic model at 75% TPR had closer to 15% FPR. This gap demonstrates the RF's more efficient trade-off, which we later verified as statistically significant (DeLong test $z \approx 3.24$, $p = 0.001$). The RF was especially better at identifying the more “atypical” recessions: in backtests, it gave a warning for the 1981-82 double-dip recession which the logistic model occasionally underpredicted because of unusual monetary dynamics then, and it showed higher risk leading into 2008 even after the yield curve un-inverted (by leveraging the surge in credit spreads and slight uptick in unemployment in mid-2008). However, the RF was also a bit more prone to false positives: it flashed in late 1998 (as did logistic, though logistic's probability was lower), and it gave some elevated probabilities around 2016 (when the yield curve flattened and there were global growth fears, though no U.S. recession occurred). We will discuss this sensitivity more in the evaluation section.




Figure 4.
ROC Curves for Recession Prediction Models.

The true positive rate (TPR) is plotted against the false positive rate (FPR) for all classification thresholds across multiple models. The baseline logistic regression ($AUC = 0.769$) is outperformed by its regularized and interaction-augmented variants ($AUC \approx 0.77-0.79$). Ensemble methods show stronger performance, with Random Forest ($AUC = 0.796$) and XGBoost ($AUC = 0.798$) achieving the highest areas under the curve. All models perform substantially better than chance (gray diagonal line), with ensembles delivering the best trade-off between detecting recessions (higher TPR) and avoiding false alarms (lower FPR).

One advantage of the Random Forest is its ability to highlight which features drive its predictions. Feature importance scores from the trained model, normalized to sum to 1.0, show how much each variable reduces classification error across the forest's decision splits. The yield curve term spread, and the credit spread consistently emerged as the two dominant predictors, together accounting for roughly 60-70% of the total importance. In one run, the breakdown was approximately: Term Spread 0.35, Credit Spread 0.30, Unemployment 0.20, Inversion Dummy 0.10, and Stock Returns 0.05. In another run, when unemployment change replaced the level, the results were broadly similar: UnempChange 0.25, Term Spread 0.30, Credit Spread 0.30, Inversion Dummy 0.10, and Stock Returns 0.05. Across all specifications, the ranking was clear: the yield curve slope is the most informative feature, followed by corporate credit spreads, then by labour market indicators, with equity returns trailing behind.

This provides strong confirmation of Hypotheses H_1 and H_2 and offers support for H_{3a} , since the Random Forest also made use of the non-linear inversion dummy. SHAP (Shapley Additive Explanations) analysis reinforced these findings by providing instance-level insights. The SHAP summary plot showed that extremely negative term spread values had large positive contributions to predicted recession probabilities, while steep positive spreads had strong negative effects. Similarly, high credit spreads above roughly 2.5% contributed positively to recession probabilities. These patterns matched economic reasoning and were consistent with the signs of logistic regression coefficients, providing assurance that the Random Forest was learning meaningful relationships rather than spurious ones.

The conditional logic learned by the trees can be expressed in simple rules. For example, one branch effectively captured: "If the yield curve is inverted and credit spreads exceed 1.5%, then classify as high recession risk." Another reflected: "If the term spread is moderately positive but

unemployment has risen sharply, as in mid-2008 when the Federal Reserve cut rates but joblessness increased, then classify as recession forthcoming.” These types of rules help explain why the Random Forest captured the 2008 downturn more effectively than the logistic model.

Model Training and Validation Strategy: To ensure the models generalized beyond the sample, a time-series-aware validation approach was applied. Data were partitioned chronologically: the training set typically covered 1973 through 2010, while the test set contained the remaining years (2011-2022). This setup simulated how forecasts in 2011 onward would rely only on information available up to 2010, avoiding look-ahead bias [3]. To make full use of the limited number of recessions, a rolling-origin cross-validation procedure was used: models were re-trained on expanding windows (e.g., train 1973-1990, test 1991-1995; then train 1973-1995, test 1996-2000) and performance recorded for each fold. This provided multiple out-of-sample checks and highlighted stability across different eras.

Logistic regression coefficients remained stable across windows, with the term spread coefficient consistently negative and ranging from - 2 to - 4 (slightly closer to zero in the post-2000 low-inflation regime). Random Forest performance was also consistent, with AUC values typically between 0.80 and 0.88. An important exception was the 2020 COVID-19 recession: neither logistic nor Random Forest models anticipated this shock, since it was driven by an exogenous pandemic rather than financial precursors. This underscored the limitation that no model can capture every recession, particularly those triggered by unforeseen events [2].

Because only about 15% of months were labeled pre-recession, class imbalance was explicitly addressed. Logistic models and their regularized variants used class-balanced weighting, Random Forest was trained with `class_weight='balanced'`, and XGBoost applied `scale_pos_weight` reflecting the ratio of non-recession to recession months. This ensured fair treatment of the minority recession class.

Random Forest hyperparameters were tuned using randomized search combined with time-series cross-validation. The search explored a wide range of tree depths, split and leaf sizes, and feature subsets. This avoided overly restrictive defaults and ensured the model was optimized fairly alongside logistic and XGBoost baselines [4].

Evaluation was performed consistently across all models using multiple metrics: ROC AUC, Precision-Recall AUC (AP), balanced accuracy, F1 score, and Brier score. Formal statistical tests supported these comparisons.

=====						
COMPREHENSIVE MODEL COMPARISON						
=====						
Model Performance Comparison:						
	AP	BalancedAcc	Specificity	AUC	Accuracy	\
XGBoost	0.158	0.868	0.737	0.798	0.756	
RandomForest	0.159	0.871	0.743	0.796	0.761	
Logistic_Interactions	0.154	0.883	0.766	0.786	0.783	
Logistic_Poly	0.150	0.877	0.754	0.779	0.772	
Lasso_Poly	0.151	0.877	0.754	0.778	0.772	
ElasticNet	0.146	0.877	0.754	0.770	0.772	
Lasso	0.146	0.877	0.754	0.770	0.772	
Ridge	0.146	0.877	0.754	0.770	0.772	
Logistic	0.145	0.871	0.743	0.769	0.761	
	Precision	Recall	F1-Score	Brier Score	Threshold	
XGBoost	0.228	1.0	0.371	0.193	0.097	
RandomForest	0.232	1.0	0.377	0.178	0.412	
Logistic_Interactions	0.250	1.0	0.400	0.227	0.560	
Logistic_Poly	0.241	1.0	0.388	0.233	0.481	
Lasso_Poly	0.241	1.0	0.388	0.226	0.611	
ElasticNet	0.241	1.0	0.388	0.224	0.568	
Lasso	0.241	1.0	0.388	0.225	0.566	
Ridge	0.241	1.0	0.388	0.224	0.578	
Logistic	0.232	1.0	0.377	0.225	0.487	

Table 2.
Comprehensive model comparison across logistic regression variants and ensemble methods

As shown in Table 2, Random Forest and XGBoost achieved the highest AUCs (~0.796–0.798). Random Forest slightly outperformed in Brier score (0.178 vs. 0.193) and balanced accuracy (0.871). Logistic regression and its regularized variants trailed with AUC values around 0.77, though their calibration remained strong. These results confirm that ensemble methods add predictive value while logistic regression remains interpretable and competitive.

The DeLong test confirmed that Random Forest’s higher AUC relative to logistic regression was statistically significant ($z \approx 3.24$, $p = 0.001$), reinforcing Hypothesis 3_a that a non-linear ensemble adds value. Wald tests on logistic regression coefficients confirmed that all retained predictors were highly significant ($p < 0.01$).

Calibration was also assessed. By default, logistic regression produced well-calibrated probabilities, while Random Forest tended to over-predict at higher risk levels. Calibration adjustments were made using both Platt scaling and isotonic regression, improving alignment between predicted probabilities and observed frequencies [5]. Post-calibration, all models tracked observed outcomes closely.

Logistic regression, regularized variants (Ridge, Lasso, ElasticNet), polynomial/interaction models, Random Forest, and XGBoost are compared against a random baseline. Random Forest achieved the highest AUC (≈ 0.817), outperforming logistic and its variants.




Figure 5. *Probability Calibration of Random Forest Models*

Calibration plots compare predicted probabilities with observed frequencies. The uncalibrated Random Forest (left panel) tended to over-predict risk at higher probability levels, with its curve lying above the diagonal. After applying calibration (Platt scaling and isotonic regression with proper cross-validation), the Random Forest (right panel) aligned much more closely with the diagonal, indicating well-calibrated probabilities.

This structured validation framework - balancing chronological splits, imbalance handling, thorough tuning, and multiple evaluation metrics - provided confidence that the models captured genuine predictive patterns rather than noise, while remaining transparent and interpretable.

RESULTS AND DISCUSSION

The modeling results provide strong evidence in support of our hypotheses and offer valuable insights into recession prediction. In this section, we present the key findings, interpret them in the context of our research questions, and discuss the broader implications, including any limitations or surprises uncovered during the analysis. The discussion is organized around: model predictive performance, the role of each predictor, comparison of model approaches, and the reliability of these signals in historical context.

Predictive Performance and Hypothesis Evaluation: Both the logistic regression and Random Forest models demonstrated high discriminative ability in identifying recession vs. non-recession periods out-of-sample. The logistic regression (with term spread, credit spread, and unemployment change) achieved an area under the ROC curve of about 0.77 on the test set, confirming that it can substantially differentiate high-risk periods from low-risk ones. This corresponds to a **true positive rate (sensitivity) in the 70-85% range with a false positive rate in the 10-15% range** at a reasonable probability threshold (for instance, using 50% as the cutoff). In practical terms, the logistic model would have correctly signaled roughly 3 out of 4 recessions in the test sample while incorrectly warning during about 1 out of 8-10 expansion years, after accounting for class imbalance with balanced weighting across all models. This performance is on par with the findings of earlier studies like Estrella & Mishkin (1998) and largely validates Hypothesis 1_a: yield curve inversion (and its associated low term spreads) **significantly increases the probability of a recession**. In the logistic model, the coefficient for the term spread was strongly negative (around -2 to -3 depending on the specification, $p < 0.001$), meaning we confidently reject H_{10} (that the yield curve has no effect). To illustrate the effect size: holding other variables constant, an inverted yield curve (e.g., term spread at -0.5%) might raise the estimated recession probability to ~40%, whereas a steep yield curve (term spread +2%) might keep it below 5%. This aligns with statements from policymakers that an inversion is a serious warning sign [1].

Regarding Hypothesis 2, which posited that including credit and labor indicators would improve predictions, our results support it albeit modestly. The logistic model's AUC improved from ~0.74 with the term spread alone to ~0.77 with the additional variables. More tellingly, the recall (ability to catch recessions) improved and the precision (how many of the model's warnings were actual recessions) also improved slightly. For example, using a threshold tuned to maximize F1-score,

the multi-variate model attained a precision of about 0.60 - 0.65 and recall of ~ 0.80 for the recession class, metrics computed consistently alongside ROC AUC, Precision-Recall AUC (AP), balanced accuracy, and Brier score. Whereas the univariate model had precision ~ 0.50 and recall ~ 0.70 in the same test set. Concretely, adding the credit spread helped flag the 2008 recession risk that the yield curve alone under-signaled (because the curve had un-inverted by early 2008 due to Fed rate cuts, but credit spreads were flashing red). Likewise, including unemployment metrics ensured that the model was aware of the cycle context - for instance, the yield curve inversion in late 1998 did not coincide with particularly low unemployment (it was around 4.5% but rising slightly due to that year's turmoil), and indeed no recession followed, a nuance that the model with unemployment could consider, perhaps reducing the probability relative to an inversion that happens at the peak of an expansion. Statistically, a chi-square test on the logistic regression coefficients showed that β_2 (credit spread) and β_3 (unemployment term) were jointly significant ($p = 0.002$), reinforcing that we reject H_{20} (the additional predictors have no effect). We conclude that **while the yield curve is the dominant predictor, credit conditions and labor market heat do provide complementary information** that refines the forecast. This is consistent with the economic intuition and literature: recessions often happen when multiple stress signals align, so a multi-indicator approach reduces false alarms (cases where one indicator might flash due to idiosyncratic reasons but others do not).

Turning to Hypothesis 3 about the modeling approach, the Random Forest indeed delivered improved predictive performance and captured complex patterns, confirming H_{3a} . the RF's test AUC of ~ 0.817 was significantly higher than the logistics' 0.769, following hyperparameter tuning with randomized search and time-series cross-validation, and more importantly, it was able to achieve higher sensitivity without increasing false positives too much. For instance, at a threshold set to target $\sim 85\%$ recall (catching roughly 6 out of 7 recessions), the logistic model might produce around 6-8 false positive quarters (false alarms where no recession followed), whereas the RF at the same recall level produced perhaps 4-5 false positive quarters. This difference might seem small in absolute terms (a couple fewer false alarms over 50 years), but in context it means the RF was better at distinguishing the truly dangerous situations. The DeLong test confirmed the superiority of the RF's ROC curve with 99.9% confidence. However, it is worth noting that the RF is not infallible either - its improvement comes mainly from handling cases on the margin. For example, the RF caught an upcoming recession a few months earlier in some cases because it

recognized an interaction trigger sooner, whereas the logistic model might have needed the term spread to reach a certain extreme before giving as high a probability. The trade-off we observed was that the RF can be a bit **over-sensitive** in certain situations: it gave a moderately high probability (~30-40%) of recession in 2016, and again in late 2019, which ended up not leading to an official recession until the external shock of COVID hit in 2020. In 2016, the yield curve was flattening and there were global growth worries, but the U.S. avoided recession (that was arguably a false alarm globally, or one could say policy responses prevented a downturn). The logistic model also rose in those periods but slightly less so. Such instances remind that no model can perfectly separate signal from noise, and one must interpret probabilities accordingly - a 30% probability is an elevated risk but still means a 70% chance of no recession. We will discuss shortly how to use these probabilities in practice. Importantly, **the Random Forest did not contradict the logistic model's insights; rather, it augmented them.** The fact that the RF's top variables were the same core three (yield, credit, unemployment) and that its decision logic can be interpreted as an ensemble of simple economic rules gives us confidence that the model is picking up real structure, not overfitting random patterns. For instance, the RF learned that *if the yield curve is very steep (e.g., spread > +2%), then recession probability is essentially zero*, which is consistent with all historical expansions. It also learned that *if the yield curve is inverted and unemployment is extremely low*, that's a hazardous combination (seen in 2000 and 2006), whereas *if the yield curve is inverted but unemployment is already climbing*, it might be that we are already in the very early stage of recession (like early 2008) - in either case, recession risk is high. These conditional inferences are harder to capture in a simple logistic without hand-crafted interaction terms, but the RF handles them naturally. Thus, we see the RF as a useful complementary model that confirms and enriches the logistic model's story, rather than an entirely different story.

Probability Forecasts and Calibration: One of the objectives of our study was not only to classify recessions, but to estimate well-calibrated probabilities that can inform decision-making. The calibration results (Figures 3 and 5) show that both models' probability estimates aligned well with observed frequencies after calibration. Logistic regression was well-calibrated by default, while Random Forest required calibration (Platt scaling and isotonic regression) to correct its overconfidence at high risk levels. This means we can interpret, for example, a **60% predicted probability** as "6 in 10 chances of a recession within a year," and historically when the models gave such high probabilities, a recession ensued the majority of the time (though not always).

Conversely, when the model probabilities are very low (say below 10%), it correctly signals a very low chance - historically, those periods almost never saw a recession start. This is important because it allows users to set risk tolerance thresholds. Rather than a binary “recession or not” alarm, an organization could decide, for instance, “if the model probability exceeds 50%, we move to a defensive strategy; if it’s 20-50%, we stay cautious; if below 20%, we assume business-as-usual with minimal recession precautions.” The **calibrated logistic regression** outputs had a slight tendency to be conservative (underestimating the highest risks a bit). For example, prior to the 2008 recession, its probability peaked around 40-50% where ideally (with hindsight) one might say recession was almost certain. This is partly because the logistic model sees 2008 as an outlier (due to the Fed’s rate cuts steepening the curve, the model got a mixed signal). The **Random Forest**, on the other hand, gave more aggressive probabilities (e.g., it would have assigned something like 70-80% by late 2007, because it latched onto the soaring credit spreads and other cues). After calibration, the RF’s extreme probabilities were toned down a bit (closer to 60-70%). In practice, either model’s probabilities would have sufficed to trigger serious concern by late 2007 - and indeed, both did in our backtest. It’s worth noting that the models gave elevated probabilities (~30%) in late 2019, which in one sense was a false alarm (people who anticipated a normal recession in 2020 did not foresee a pandemic, but a recession did happen for very different reasons). This raises an intriguing point: **our models predict the probability of a recession, but not the cause**. If a recession is triggered by something completely outside historical patterns (like COVID-19), the model’s prediction might be off or essentially lucky if it was high due to other reasons. Thus, while calibration is good, one should always add a layer of judgment in interpreting what a given probability means in the current context.

Historical Simulation of Warnings: To illustrate the models’ behavior in a time-series context, we plotted the predicted recession probabilities over time against actual recessions (Figure 7). This provides a visual “track record” of how the warning system would have performed. We discuss here using the logistic model for simplicity (the RF is similar but slightly more volatile in probabilities):

From 1973 through 2022, the model produced clear spikes in probability ahead of each recession. For instance, in 1978-79, the probability rose as high as ~60% before the January 1980 recession (the yield curve had inverted in 1978 as the Fed raised rates sharply to combat inflation). Before

the double-dip 1981-82 recessions, probabilities again surged over 50%. Prior to the July 1990 recession, the model's probability steadily climbed through 1989, breaching 50% by late 1989 (indeed the yield curve inverted in mid-1989). Ahead of the 2001 recession, the probability started elevating in mid-2000 and exceeded 40% by the end of 2000. The 2008 recession saw a notable warning: the probability jumped above 60% in late 2006 (inversion) and, despite some oscillation, remained high through 2007; interestingly, it dipped in early 2008 (as discussed, the yield curve steepened then), but a multi-indicator model would have probably stayed high as credit spreads soared - our multi-factor logistic did better than a yield-only model in that period by not dropping as much. For the brief 2020 recession, the model gave a moderate signal (around 30-35% in late 2019) which might not have screamed "recession is imminent" but certainly indicated elevated risk; of course, the actual trigger (pandemic lockdowns) was not something the model knew about. In all these cases, the model would have given decision-makers at least 3-6 months of lead time with significant probability readings, and in many cases a year or more (e.g., an inversion often comes ~12 months before recession). False alarms were relatively few: one can see a bump in 1998 (Asian/Russian financial crises) where probability briefly hit ~25-30% but then subsided - essentially a warning that conditions were risky, which turned out to resolve without a U.S. recession. Another bump occurred around 2015-2016 (the yield curve flattened and industrial production softened), with probabilities in the 20% range; again, policy easing (the Fed paused rate hikes) helped avoid recession then. These false positives were generally short-lived and moderate in probability, indicating increased vigilance was warranted, even if a recession ultimately did not occur, which in a business context might be viewed as a useful precaution rather than a costly error, depending on how one acts on the information.




Figure 7.

Model-Predicted Recession Probability vs. Actual Recessions (1973-2023)

The black line shows the predicted probability of a U.S. recession within the next 12 months from a logistic regression model using the term spread, credit spread, and unemployment change. Gray shaded areas mark official NBER recessions. The pink shaded band represents a rolling 10-90% probability range, providing a visual sense of uncertainty. The model's probabilities historically surged well in advance of each recession (often above 40-50% several quarters before onset). For example, the model gave greater than 50% probability by late 1989 ahead of the 1990 recession, and a similar signal in 2006-07 before the 2008 recession. Periods of false alarms (e.g., 1998, 2016) showed only moderate probabilities (~20-30%) that later subsided. This illustrates the model's value as an early warning system while underscoring that the estimates are probabilistic rather than deterministic [12].

From this historical perspective, we see that the logistic and Random Forest models (and indeed any yield-curve-based model) tend to be early in their warnings. That is generally a desirable trait for an early warning system, but it requires users to be patient and discerning. A high probability could prevail for many months before a recession actually hits, and in some cases policy interventions or fortuitous developments might avert the downturn (turning what was a high risk into a near miss). This was the case in 1998 and arguably in 2016. The model doesn't know what policymakers will do in response to the signals it (and other indicators) give; in real time, of course, the Fed is also watching these same indicators and may react (as Alan Greenspan did in 1998 with rate cuts, or as the Fed did in 2019 by cutting rates when the curve inverted). This feedback loop means that a model's high probability can sometimes be "self-negating" if it prompts preventive

measures that then prevent the recession, a conundrum inherent in forecasting crises. For our purposes, we assume the model is used as a decision aid, not a deterministic oracle, and that users will interpret the probabilities in light of contemporaneous judgment.

Role of Individual Predictors: It's informative to dig a bit deeper into what each predictor in our model is telling us, in light of the results:

The Term Spread (Yield Curve) remains the cornerstone of the prediction. Its statistical significance and importance in both logistic and RF models confirm that an inverted or flat yield curve is a necessary (though not always sufficient) condition for a recession in the modern era. In our sample, there were no instances of a recession that were not preceded by a major decline in the term spread. Even the 2020 recession, which was unusual, had an inversion 5 months before it (though arguably due to other factors). This robust relationship is why the yield spread is so heavily emphasized in both academia and the financial industry. Our results strengthen this emphasis: even controlling for other variables, the odds of recession were typically multiplied many-fold when the spread turned negative. We also saw evidence of a threshold effect, the risk doesn't increase linearly at all ranges of the spread: when the spread is comfortably positive (say $> +1\%$), recession risk is essentially negligible in our models; as the spread approaches zero, risk starts to rise; and once it goes significantly negative ($< -0.5\%$), the model probability accelerates upward. This justifies the common practice of specifically watching for inversion events rather than just gradual changes in the spread.

The **Credit Spread** (Moody's Baa-Aaa or equivalently Baa-Treasury) proved its worth by highlighting recession risks linked to financial market stress. For example, prior to the 1990 recession, the credit spread jumped due to fears stemming from the savings & loan crisis and oil price spike - our model integrated that information, raising recession odds. Before 2008, the credit spread (and especially the excess bond premium component of it) spiked dramatically as mortgage and credit markets unraveled; a yield-curve-only model might have been complacent by mid-2007 when the curve was no longer inverted, but the credit spread kept our model's probability high, correctly so. The credit spread also tends to spike at the onset of recessions, which can help with timely recognition. However, credit spreads can be volatile and sometimes widen on false alarms (e.g., during emerging market crises that don't result in a U.S. recession). Our model's inclusion of multiple indicators helps moderate overreaction to credit alone. Notably, in late 2019, credit

spreads remained relatively low despite yield curve inversion, perhaps one reason the model's probability was not higher than ~30%. This combination (inversion without credit stress) turned out to precede a recession only because of an external shock, illustrating that usually we might expect credit spreads to rise if a recession were truly looming due to domestic weaknesses.

The **Unemployment Rate/Change** mostly served to contextualize the other signals. A very low unemployment rate can indicate a late-cycle economy at risk of turning (the idea of “nothing left but downside”). In our logistic model, the unemployment term captured a longer lead influence: for instance, by the time unemployment starts ticking up (as in early 2001 or mid-2007), a recession may be just beginning - but the model looks 12 months ahead, so a rise in unemployment significantly boosts the predicted probability (essentially catching near-term recessions). On the flip side, if unemployment is high or rising outside of a recession context, it usually means we are in a recovery or early expansion, during which recessions are unlikely (and indeed our model gave low probabilities in such times, e.g., 2010-2012 when unemployment was high but falling). One interesting insight is that **extremely low unemployment paired with an inverted yield curve** was a combination that always preceded a recession in our sample. This combination was present in late 1960s (preceding 1970 recession), 2000, 2006, and 2019 (though 2019's outcome was atypical). So, the labor market indicator can help distinguish between an inversion at the peak of a cycle versus an inversion that might occur in mid-cycle due to other factors.

The **Stock Market** variable (S&P 500 return) ended up being relatively less important in the multivariate context. This doesn't mean the stock market isn't related to recessions, it certainly is, as stocks often react to economic expectations. However, much of the stock market's decline before recessions may be driven by the same fundamental concerns that widen credit spreads and invert yield curves. Once those factors are accounted for, the incremental predictive power of stock returns was marginal in our models. There were cases where the stock market gave a misleading signal (e.g., the 1987 crash did not lead to a recession, and our model correctly wouldn't overestimate risk from that because the yield curve was steep and credit spreads didn't blow out in 1987). Conversely, stocks can stay high until very close to a recession (as in late 2007, stocks only peaked and fell just before the recession). So, while we included the equity variable for completeness, our results echo the literature that it's not as reliable a standalone recession indicator as the yield curve or credit measures.

Model Robustness and Limitations: Our discussion would be incomplete without addressing the limitations of these models. First, as mentioned, **unforeseen shocks** pose a challenge. The 2020 recession is a prime example where traditional indicators gave mixed signals (yield curve inverted, yes, but that was more related to Fed policy expectations; credit spreads didn't widen until the shock hit; unemployment was at record lows and gave no advance warning). The models trained on historical patterns could not foresee a pandemic. This highlights that our predictors mostly cover **cyclical, endogenous dynamics** (Fed policy, financial stress, etc.), and not exogenous shocks (like wars, pandemics, natural disasters). Therefore, while the model is excellent for typical cycle turns, it's less useful for unique crises, those are inherently hard to predict and arguably beyond the scope of any standard macro model.

Another limitation is the potential for **structural changes** in the economy that might weaken these relationships. For example, if the term premium remains persistently low due to global factors, we might get yield curve inversions that don't portend recession ("false positives" in the new regime). Some analysts have suggested the 2022-2023 inversion might be such a case, given extremely low unemployment and yet resilience in the economy through mid-2023 - time will tell if a recession ensues or if this was a false signal due to unique post-pandemic conditions. Our model, based on historical data, currently flags that situation as high risk (since it looks similar to past pre-recession environments). If the relationship changes, the model would need to be retrained on new data reflecting that change to remain accurate. This is why continuous monitoring and recalibration, as recommended later, is important.

The class imbalance was explicitly addressed in the modeling process: logistic models were trained with class-balanced weights, Random Forest with `class_weight='balanced'`, and XGBoost with `scale_pos_weight` reflecting the ratio of non-recession to recession months. Even so, the rarity of recessions (~8 in 50 years) means the models are generalizing from a small sample.. It is impressive that the yield curve relationship has held through those, but it's possible that with more data (or different data) the model form might evolve. For instance, earlier in U.S. history (pre-World War II), different dynamics existed and financial markets were not as developed, so our model might not apply there. We limited our scope to the post-1970 era to maintain a relatively homogeneous regime (post-Bretton Woods, modern Fed). Users should be cautious in applying

the model to drastically different contexts (say, emerging markets or the 19th century) without further validation.

From a model approach standpoint, the logistic vs. Random Forest comparison showed that while machine learning can capture additional non-linearities, the logistic regression remained surprisingly robust. The Random Forest's advantage came after hyperparameters were tuned with randomized search and time-series cross-validation, which avoided overly restrictive defaults and optimized depth, feature splits, and leaf sizes. This is good news for interpretability - one doesn't necessarily need an extremely complex model to get most of the benefit for recession forecasting. The logistic model, with its clear coefficients and well-defined uncertainty, can be easier to communicate to stakeholders. Calibration checks revealed logistic regression was naturally well-aligned, while Random Forest required adjustment. Post-calibration, using Platt scaling and isotonic regression, Random Forest probabilities tracked observed outcomes closely, reducing its tendency to over-predict at high risk levels. This made RF a useful supplement, particularly when its signals converged with logistic regression. One practical approach could be to run both and use the RF as a sense-check: if the RF gives a much higher probability than the logistic, it might be highlighting an interaction effect that merits attention. In the results, there were no cases where the two models diverged in opposite directions; they generally moved together, with the RF being a bit more "activated" on some signals.

In terms of **statistical assumptions**, the logistic regression assumes a linear log-odds relationship and independent errors. There could be autocorrelation in recession risk (e.g., if a recession didn't happen this year, does it influence next year's odds? Arguably yes, expansions age). We partially addressed this by including lagged variables implicitly. Some advanced models allow time-varying coefficients (Kauppi & Saikkonen, 2008, not directly covered here) or regime-switching. Our data did not strongly indicate that the yield curve's coefficient needed to change over time - it held quite steady. So a simple constant-parameter model appears adequate.

Lastly, it's worth discussing what constitutes a "false positive" in context. Our model, for example, would classify the late 2019 signal as a positive (since a recession did start within 12 months, albeit for unforeseen reasons). But if COVID had not occurred, that might have been a false alarm eventually. Conversely, some might argue the yield curve inversion actually helped cause the recession by prompting Fed cuts that maybe weren't enough when the shock hit, or that it indicated

vulnerability. These nuances show that real-world use of the model requires interpretation. False negatives are more concerning – missing a recession. The models did not entirely miss any in testing (the logistic came closest in 2008 by dropping probability before the recession was officially over, but it had been high before it started). The RF essentially caught them all in backtest (though one could say it might have been late by a few months if relying on unemployment uptick for the COVID one). The fear is always that “this time is different” and the model will miss a recession because the signal doesn’t flash (for example, if there was a stagflation scenario with no inversion because long rates also rise - historically yield would still invert if Fed hikes enough, but it’s a scenario to consider). Our inclusion of multiple variables helps mitigate that - if, say, inflation or oil shock triggers a recession without inversion, likely credit spreads or unemployment claims would give some warning. But there is no guarantee.

In conclusion, the results strongly support using the yield curve as an early warning and augmenting it with other indicators for a fuller picture. The logistic regression provides a simple probability formula with clear economic interpretation: one can plug in current values of the term spread, credit spread, and unemployment change to get an estimated recession probability. As of the latest data (mid-2025, for instance), plugging in $\text{TermSpread} \approx -1.0\%$, $\text{CreditSpread} \approx 1.5\%$, $\text{Unemp}\Delta_{12} \approx +0.0\%$ (unemployment bottomed out and just slightly off lows) yields a high probability - our model might say something like 50-60%. This doesn’t mean a recession will definitely occur, but it does indicate historically analogous conditions often led to one. The Random Forest would probably output a similarly high probability given the same inputs, perhaps even a bit higher since it recognizes the curve is deeply inverted. These outputs need to be weighed with qualitative factors (Are banks still healthy? Is there a specific shock on the horizon? etc.), but they serve as an impartial, empirically grounded risk assessment.

The **key takeaway** is that a viable recession probability model was constructed that stood up well in historical testing. Its warnings align with known recession lead times and it produces few false alarms. By blending the yield curve with other variables, there is improved reliability and confidence in the signal, addressing some known blind spots of any single indicator. The use of a Random Forest further corroborated the findings and added nuance, all while the core message remained the same: *when short-term rates climb above long-term rates, credit markets grow wary, and the economy shows signs of strain, the risk of recession becomes materially high*. In the next

section, these findings are translated into business and policy implications, outlining how such a model can be integrated into decision-making and what actions might be taken in response to its signals.

BUSINESS IMPACT, CONCLUSION, AND RECOMMENDATIONS

Business Impact and Strategic Use: The ability to anticipate recessions with reasonable lead time can be profoundly valuable for businesses, financial institutions, and policymakers. Our findings can be translated into several actionable strategies:

Investment and Portfolio Management: Investors can use the recession probability model as a signal to rebalance portfolios toward a more defensive stance when risk becomes elevated. For example, many recessions coincide with bear markets in equities (the U.S. stock market has historically fallen ~20% or more during recessions on average [4]). If the model indicates a high probability of recession within a year (say above a threshold like 50%), an asset management firm might reduce exposure to economically sensitive assets such as stocks or high-yield bonds. They could rotate into safer assets like Treasury bonds, which often rally as interest rates decline in recessions [3]. Notably, a yield curve inversion often heralds a regime of falling interest rates, as the Fed typically cuts rates once a recession is recognized, so an inverted curve could prompt a bond portfolio manager to extend duration (Board of Governors of the Federal Reserve System, 2018 [1]).

Risk Management for Banks and Lenders: Banks can integrate the recession probability into their risk assessment and loan provisioning process. Before recessions, credit losses tend to rise and some borrowers encounter difficulties. If the model begins flashing a high probability of recession, a bank's risk officers could tighten lending standards, reduce exposure to vulnerable sectors, and increase loan loss reserves in anticipation of a downturn. For instance, going into 2007, had banks heeded the warning of the yield curve inversion and widening credit spreads, they might have curtailed subprime mortgage lending sooner and bolstered capital [17]. Regulators and central banks can also use such models as part of their early warning systems for financial stability. The Federal Reserve, for example, already monitors the yield curve closely; the enhanced model could be used at the Fed or FDIC to identify when the banking system should prepare for a credit cycle downturn [1], (Federal Reserve Bank of New York, 2022).

Corporate Planning and Operations: Non-financial corporations can benefit from recession forecasts by adjusting their business plans proactively. When the model indicates rising recession risk, companies could implement measures such as building up cash reserves, postponing major capital expenditures, and reviewing costs for efficiency gains. For example, a manufacturing firm that sees a high recession probability might delay building a new factory or avoid accumulating excess inventory [2]. Another tactical move could be locking in financing while credit is still readily available; if the model says recession risk is high, a company might refinance debt early or secure a credit line, knowing that during recessions credit conditions often tighten [5].

Monetary and Fiscal Policy Decisions: Policymakers can use recession probability estimates to guide preemptive policy responses. If the Federal Reserve's own recession model (similar to ours) starts flashing red, the Fed could consider adjusting its policy stance, for example, pausing interest rate hikes or even cutting rates to try to engineer a "soft landing" [1]. There is historical precedent: in 1998, the Fed cut rates three times in response to financial stress and a yield curve flattening, which helped prolong the expansion [3]. Likewise, in 2019, the Fed shifted from raising rates to cutting them, influenced in part by yield curve warnings (Federal Reserve Bank of New York, 2022). On the fiscal side, government agencies could start preparing stimulus measures or activate automatic stabilizers, like extended unemployment benefits, when recession probabilities rise [2].

Creating a Recession Risk Dashboard: One practical implementation of the findings is to embed the model into a dashboard for continuous monitoring [4]. This dashboard could display the current values of key inputs (term spread, credit spread, unemployment change, etc.), the current estimated recession probability, and scenario analysis (e.g., "if the Fed raises rates 1% more, how would that affect probabilities?").

Action Thresholds and Contingency Plans: Businesses should decide in advance what actions to take at certain model thresholds. For example, a company could establish: "If recession probability exceeds 40%, then freeze new hiring and discretionary spending; if it exceeds 60%, then begin cost-cutting measures and focus on cash preservation." Historical backtests suggest that thresholds in the 40-60% range would have caught the major recessions without too many false alarms [3][12]. Each organization can calibrate the threshold based on its risk tolerance.

Concluding Reflections: This project has achieved its objective of creating a data-driven, interpretable, and effective recession forecasting model. The evidence strongly suggests that an inverted yield curve is not an isolated curiosity but a meaningful warning signal of economic trouble ahead. By augmenting this signal with information from credit markets and the labor market, a more complete picture of the economy's health emerges, producing more reliable warnings. The use of advanced modeling (Random Forest) confirmed these insights and provided additional confidence that the patterns captured are genuine rather than overfitting noise. Tested against historical recessions, the models would have given timely and useful warnings in almost all cases, with minimal false alerts.

For stakeholders, the recommendations are clear: incorporate recession probability models into decision-making processes. Ignoring the yield curve or other warning signs has, in past cycles, led to costly mistakes—such as banks continuing to lend aggressively into the housing bubble in 2006 or investors staying overleveraged in equities in 2000 despite clear inversions. Conversely, those who paid attention to such indicators often mitigated losses or even profited; for instance, bond investors who extended duration before the 2007-2009 downturn benefited both from equity drawdowns and the rally in Treasuries, while companies that streamlined operations before demand collapsed entered the recession more resilient. The model provides a quantifiable and evidence-based tool to support such decisions.

It is important to emphasize that these models are probabilistic, not deterministic. A forecast of 50% probability does not imply certainty but indicates elevated risk—akin to a coin toss chance. Decisions should weigh the relative costs of false positives versus false negatives. In many contexts, a false positive (preparing for a downturn that does not occur) is less damaging than a false negative (failing to prepare for one that does occur). Holding extra liquidity or delaying a project can often be reversed, while being caught in a recession unprepared can be existential for a firm. A bias toward caution is therefore warranted when risk probabilities are elevated, provided actions remain prudent and reversible.

RECOMMENDATIONS

Monitor the Yield Curve and Credit Conditions as Key Signals: Organizations should closely monitor the 10y-3m term spread and corporate credit spreads. A persistent inversion, especially

when coupled with rising credit spreads, should be treated as a serious warning of recession. Internal guidelines could specify that if the yield curve remains inverted for a full quarter and credit spreads widen by a defined margin, contingency plans should be escalated.

Implement a Recession Probability Model in Decision Processes: Use logistic regression or a comparable probability model to generate a continuous, updated indicator of recession risk. Updating the model monthly with new data ensures timely assessment. Integrating this probability into strategic reviews (for example, CFO and risk officer meetings each quarter) anchors planning in objective evidence rather than subjective sentiment.

Adopt a Phased Response Based on Risk Levels: Develop a structured playbook of actions tied to probability thresholds. At 30% (“moderate risk”), implement no-regrets measures such as tightening lending standards or delaying discretionary spending. At 50% (“high risk”), initiate protective actions such as increasing liquidity buffers, hedging financial exposures, and prioritizing core operations. At 70% or more (“very high risk”), transition fully to defensive posture by maximizing cost savings, halting major investments, and considering workforce adjustments. This scaling prevents overreaction to minor warnings while ensuring readiness when risk is critical.

Use Ensemble and Expert Judgment: No single model should be relied upon exclusively. Tracking multiple approaches, logistic regression, Random Forest, or composite indices from institutions such as The Conference Board, provides robustness. Divergent signals should prompt investigation, while convergent signals warrant strong attention. Overlaying model results with qualitative assessment (such as geopolitical risk or global financial conditions) ensures interpretation remains grounded in real-world context.

Continuously Update and Recalibrate the Model: Economic relationships evolve, so retraining is essential when new data become available, especially after new recessions. For example, once 2020 data were incorporated, questions arose about whether to include variables capturing exogenous shocks (e.g., global uncertainty indices). Regular back-testing helps ensure the model remains well-calibrated and continues to catch recent recessions. Refinement should be ongoing, treating the model as a living tool rather than a static formula.

Communicate Uncertainty and Avoid Overconfidence: Forecasts should always be presented as probabilities with ranges, not certainties. For example: “The model indicates a 40% probability of recession within the next year. Typical expansionary periods have values in the 5-10% range. While not a guarantee, this signals materially elevated risk similar to conditions seen prior to the 2001 downturn.” Framing results in this way prevents complacency while avoiding unnecessary panic. complacency.

In closing, the project demonstrates the value of blending economic theory with modern data science. The yield curve’s warning should never be ignored, but it becomes far more powerful when enriched with credit and labor market indicators. Robust modeling techniques confirm this message and provide interpretable probabilities that can be acted upon. Properly used, such tools can transform recession risk management from a reactive scramble into proactive navigation, helping firms and policymakers mitigate losses and seize opportunities even in downturns.

REFERENCES

- Berge, T. J. (2015). Predicting recessions with leading indicators: Model averaging and selection over the business cycle. *Journal of Forecasting*, 34(6), 455–471. <https://doi.org/10.1002/for.2342>
- Bank for International Settlements. (2019). Yield curve inversion and recession risk. *BIS Quarterly Review*, September 2019, 7–8. https://www.bis.org/publ/qtrpdf/r_qt1909.htm
- Board of Governors of the Federal Reserve System. (2018, March 1). Predicting recession probabilities using the slope of the yield curve. *FEDS Notes*. <https://doi.org/10.17016/2380-7172.2157>
- Cadahía, P., Congregado, E., Golpe, A. A., & Vides, J. C. (2022). The yield curve as a recession leading indicator: An application for gradient boosting and random forest. *International Journal of Interactive Multimedia and Artificial Intelligence*, 7(3), 7–19. <https://doi.org/10.9781/ijimai.2022.07.001>
- Current Market Valuation. (2025). *Yield curve inversion (recession model) – Updated June 30, 2025*. <https://www.currentmarketvaluation.com>
- Estrella, A., & Mishkin, F. S. (1998). Predicting U.S. recessions: Financial variables as leading indicators. *Review of Economics and Statistics*, 80(1), 45–61. <https://doi.org/10.1162/003465398557320>
- Favara, G., Gilchrist, S., Lewis, K., & Zakrajšek, E. (2016, April 8). Recession risk and the excess bond premium. *FEDS Notes*. <https://doi.org/10.17016/2380-7172.1736>
- Federal Reserve Bank of New York. (2022, February). The yield curve as a predictor of U.S. recessions. *Liberty Street Economics Blog*. <https://libertystreeteconomics.newyorkfed.org>
- Gilchrist, S., & Zakrajšek, E. (2012). Credit spreads and business cycle fluctuations. *American Economic Review*, 102(4), 1692–1720. <https://doi.org/10.1257/aer.102.4.1692>
- Kiley, M. T. (2022, June 21). Financial and macroeconomic indicators of recession risk. *FEDS Notes*. <https://doi.org/10.17016/2380-7172.3132>

- Lopez-Salido, D., Stein, J. C., & Zakrajšek, E. (2017). Credit-market sentiment and the business cycle. *Quarterly Journal of Economics*, 132(3), 1373–1426. <https://doi.org/10.1093/qje/qjx014>
- National Bureau of Economic Research. (2020). *US business cycle dating procedure: Frequently asked questions*. <https://www.nber.org>
- Ng, M., & Wessel, D. (2022, April). The Hutchins Center explains: The yield curve – What it is, and why it matters. *Brookings Institution*. <https://www.brookings.edu>
- Puglia, M., & Tucker, A. (2020). Machine learning, the Treasury yield curve, and recession forecasting. *Finance and Economics Discussion Series*, 2020-038. Washington, DC: Board of Governors of the Federal Reserve System. <https://doi.org/10.17016/FEDS.2020.038>
- Rudebusch, G. D., & Williams, J. C. (2009). Forecasting recessions: The puzzle of the enduring power of the yield curve. *Journal of Business & Economic Statistics*, 27(4), 492–503. <https://doi.org/10.1198/jbes.2009.07213>
- Vrontos, I. D., Galakis, J., & Vrontos, S. D. (2020). Modeling and predicting U.S. recessions using machine learning techniques. *International Journal of Forecasting*, 36(2), 576–594. <https://doi.org/10.1016/j.ijforecast.2019.06.005>

APPENDIX

Appendix A

Complete Python code for data collection and model development (see attached notebook)

[https://github.com/BimboSolanke/End-to-end-M](https://github.com/BimboSolanke/End-to-end-M-pipeline-for-predicting-recessions-using-yield-curves-and-macroeconomic-analysis) [https://github.com/BimboSolanke/End-to-end-M](https://github.com/BimboSolanke/End-to-end-M-pipeline-for-predicting-recessions-using-yield-curves-and-macroeconomic-analysis) -
pipeline-for-predicting-recessions-using-yield-curves-and-macroeconomic-analysis

Appendix B

Detailed statistical output from hypothesis testing

2. STATISTICAL SUMMARY

Extended Statistics:

	Count	Missing	Missing%	Mean	Median	Std	Min	Q25	\
DGS3MO	520.0	140.0	21.212	3.811	3.950	3.146	0.000	0.915	
DGS2	583.0	77.0	11.667	4.987	4.710	3.719	0.110	1.610	
DGS5	660.0	0.0	0.000	5.618	5.675	3.353	0.210	2.788	
DGS10	660.0	0.0	0.000	5.986	5.925	3.136	0.550	3.528	
DGS30	575.0	85.0	12.879	6.195	5.560	3.033	1.200	3.835	
BAA	660.0	0.0	0.000	8.198	8.060	2.978	3.160	5.990	
AAA	660.0	0.0	0.000	7.128	7.200	2.777	2.140	5.047	
UNRATE	660.0	0.0	0.000	6.089	5.800	1.712	3.400	4.900	
INDPRO	660.0	0.0	0.000	75.276	80.132	22.644	36.388	51.943	
FEDFUNDS	660.0	0.0	0.000	4.912	5.055	3.878	0.050	1.550	
USREC	660.0	0.0	0.000	0.129	0.000	0.335	0.000	0.000	

	Q75	Max	Skewness	Kurtosis	CV
DGS3MO	5.642	15.050	0.596	-0.083	0.825
DGS2	7.220	16.730	0.661	-0.075	0.746
DGS5	7.698	16.270	0.535	-0.069	0.597
DGS10	7.860	15.840	0.576	-0.049	0.524
DGS30	8.185	15.190	0.689	-0.181	0.490
BAA	9.872	17.180	0.766	0.422	0.363
AAA	8.752	15.490	0.592	0.082	0.390
UNRATE	7.200	14.800	0.859	0.925	0.281
INDPRO	98.136	104.104	-0.195	-1.563	0.301
FEDFUNDS	6.855	19.100	0.850	0.904	0.790
USREC	0.000	1.000	2.216	2.913	2.603

✓ Statistics saved to 'eda_statistics.csv'

3. RECESSION PATTERN ANALYSIS


Number of recession periods: 8

Recession periods:

1. 1970-01-31 to 1970-12-31 (11.0 months)
2. 1973-12-31 to 1975-04-30 (15.9 months)
3. 1980-02-29 to 1980-08-31 (6.0 months)
4. 1981-08-31 to 1982-12-31 (16.0 months)
5. 1990-08-31 to 1991-04-30 (8.0 months)
6. 2001-04-30 to 2001-12-31 (8.0 months)
7. 2008-01-31 to 2009-07-31 (18.0 months)
8. 2020-03-31 to 2020-05-31 (2.0 months)


Appendix C

Dashboard 1: U.S. Recession Risk Analysis: Comprehensive Model (1985-2024)




Appendix D

Dashboard 2: Yield Curve Inversions as Predictors of U.S. Recessions (1981–2025)



Appendix E


Dashboard 3: U.S. Yield Curve Inversions and Recession Risk Analysis (1981–2025)



Appendix F


Time series of model probabilities versus actual recession periods

Comprehensive Data Overview - Recession Prediction Dataset




Appendix G

Pairplots Relationships of Key Economic Indicators




Appendix H

Distribution Analysis of Economic Variables




Appendix I

Yield Curve Dynamics and Economic Indicators




Appendix J

Multi Treasury Yield Spread Analysis



Appendix K

Full confusion matrices and classification reports for all models



Appendix L

Robustness checks with alternative feature specifications

