

INVERSION TO CONTRACTION: QUANTIFYING U.S. RECESSION PROBABILITIES WITH MACRO-FINANCIAL SIGNALS

YIELD CURVES, CREDIT SPREADS, AND UNEMPLOYMENT IN A BLENDED ML FRAMEWORK



Instructor: Prof. Omid Isfahanialamdari

Author: Bimbo Olukoya Solanke


Course: DAMO-611-6 Data Analytics Case Study

Date: September 2025

14/09/2025

THE PROBLEM WHY THIS MATTERS



- Economic Impact: Recessions affect millions
- Current Challenge: Limited lead time for preparation
- **Key Question:** Can we reliably predict recessions 12-18 months ahead?



THE PROBLEM - WHY THIS MATTERS

Yield Curve Inversions as Predictors of U.S. Recessions (1981–2025)

Current Status	Historical Mean	% Days Inverted
● Normal	1.56 Mean Spread (%)	20.57% Trading Days Inverted



RESEARCH QUESTIONS & HYPOTHESES

H_1 : Yield curve inversions significantly increase recession probability ($p < 0.05$)

FORMAL HYPOTHESIS TESTING

HYPOTHESIS 1: YIELD CURVE INVERSION

H0: Yield curve inversion does NOT increase recession probability

H1: Yield curve inversion SIGNIFICANTLY INCREASES recession probability

TS_10y3m:

Mean (Recession): 0.237
Mean (Normal): 0.090
T-statistic: 3.831
P-value: 0.0001
Logit coefficient: 1.147
Logit p-value: 0.0003
*** REJECT H0: TS_10y3m inversion DOES predict recession ($p < 0.05$)

TS_10y2y:

Mean (Recession): 0.163
Mean (Normal): 0.105
T-statistic: 1.474
P-value: 0.1411
Logit coefficient: 0.503
Logit p-value: 0.1438
FAIL TO REJECT H0: No significant effect

TS_5y3m:

Mean (Recession): 0.312
Mean (Normal): 0.098
T-statistic: 5.303
P-value: 0.0000
Logit coefficient: 1.437
Logit p-value: 0.0000
*** REJECT H0: TS_5y3m inversion DOES predict recession ($p < 0.05$)

TS_5y2y:

Mean (Recession): 0.312
Mean (Normal): 0.128
T-statistic: 4.205
P-value: 0.0000
Logit coefficient: 1.135
Logit p-value: 0.0001
*** REJECT H0: TS_5y2y inversion DOES predict recession ($p < 0.05$)

TS_2y3m:

Mean (Recession): 0.350
Mean (Normal): 0.100
T-statistic: 6.062
P-value: 0.0000
Logit coefficient: 1.578
Logit p-value: 0.0000
*** REJECT H0: TS_2y3m inversion DOES predict recession ($p < 0.05$)

TS_30y10y:

Mean (Recession): 0.138
Mean (Normal): 0.037
T-statistic: 3.650
P-value: 0.0003
Logit coefficient: 1.409
Logit p-value: 0.0007
*** REJECT H0: TS_30y10y inversion DOES predict recession ($p < 0.05$)

TS_30y5y:

Mean (Recession): 0.125
Mean (Normal): 0.050
T-statistic: 2.542
P-value: 0.0114
Logit coefficient: 0.999
Logit p-value: 0.0145
*** REJECT H0: TS_30y5y inversion DOES predict recession ($p < 0.05$)

RESEARCH QUESTIONS & HYPOTHESES

HYPOTHESIS 2: CREDIT SPREAD

H0: Credit spread widening does NOT affect recession probability

H1: Credit spread increase SIGNIFICANTLY RAISES recession probability

Credit Spread (BAA-AAA):

Mean (Recession): 1.183

Mean (Normal): 0.936

T-statistic: 5.788

P-value: 0.0000

Logit coefficient: 1.546

Logit p-value: 0.0000

*** REJECT H0: Credit spread DOES predict recession ($p < 0.05$)

HYPOTHESIS 3: UNEMPLOYMENT CHANGE

H0: Rise in unemployment does NOT increase recession odds

H1: Unemployment uptick SIGNIFICANTLY INCREASES recession probability

Unemp_D12:

Mean (Recession): 0.495

Mean (Normal): -0.203

T-statistic: 3.984

P-value: 0.0001

Logit coefficient: 0.292

Logit p-value: 0.0004

*** REJECT H0: Unemp_D12 DOES predict recession ($p < 0.05$)

Sahm:

Mean (Recession): 0.350

Mean (Normal): 0.135

T-statistic: 4.764

P-value: 0.0000

Logit coefficient: 1.238

Logit p-value: 0.0000

*** REJECT H0: Sahm DOES predict recession ($p < 0.05$)

THE COMPLETE DATASET - 54 YEARS OF ECONOMIC INDICATORS



Source	Variables Collected	Purpose in Dataset
FRED – Federal Reserve Economic Data (St. Louis Fed)	Treasury Constant Maturity Yields (3M, 2Y, 5Y, 10Y, 30Y); Moody’s Seasoned Baa and Aaa Corporate Bond Yields; Effective Federal Funds Rate	Construct yield curve spreads (e.g., 10y–3m); compute credit spreads; capture monetary policy stance
BLS – Bureau of Labor Statistics	Civilian Unemployment Rate (monthly, seasonally adjusted)	Measure labor market tightness and cycle turning points
NBER – National Bureau of Economic Research	Official recession dates (monthly chronology)	Define the target variable (recession vs. non-recession months)
Yahoo Finance / Bloomberg	S&P 500 Index (monthly closes)	Compute 12-month stock market returns as a complementary predictor
CBO – Congressional Budget Office	Natural rate of unemployment (NAIRU)	Calculate unemployment gaps (overheating vs. slack conditions)

DATA COLLECTION SUMMARY


Total observations: 660
Date range: 1970-01-31 to 2024-12-31
Number of features: 11
Missing values per column:
- DGS3MO: 140 (21.2%)
- DGS2: 77 (11.7%)
- DGS30: 85 (12.9%)

Recession Statistics:
- Recession months: 85 (12.9%)
- Non-recession months: 575 (87.1%)



Spread	Mean	Std	Min	Max	% Inverted	Correlation at 19-month lead
10Y-3M	1.57	1.24	-1.88	4.09	8.33%	-0.354
10Y-2Y	0.86	0.92	-2.01	2.84	14.24%	-0.244
5Y-3M	1.12	1.00	-1.78	3.93	9.70%	-0.303
5Y-2Y	0.45	0.54	-1.32	1.58	17.12%	-0.167
30Y-10Y	0.37	0.41	-0.71	1.37	11.82%	-0.177
30Y-5Y	0.78	0.80	-1.33	2.82	10.91%	-0.163
2Y-3M	0.62	0.66	-1.30	3.35	10.30%	-0.120

Treasury Spread & Prediction Windows

D. Lead-Lag Correlation: Yield Spread vs Future Recession



WHY FOCUS ON THE 10-YEAR – 3-MONTH SPREAD?



1. Historical Benchmark
2. Economic Meaning
3. Stability Over Time
4. Clarity & Communication

All Models Performance Metrics

F11	AUC	Accuracy	Brier Sc..	F1-Score	Precision	AP	Balance..	Recall	Thresho..
ElasticNet	0.770	0.772	0.224	0.388	0.241	0.1459	0.8772	1.000	0.5679
Lasso	0.770	0.772	0.225	0.388	0.241	0.1459	0.8772	1.000	0.5663
Lasso_Poly	0.778	0.772	0.226	0.388	0.241	0.1509	0.8772	1.000	0.6107
Logistic	0.769	0.761	0.225	0.377	0.232	0.1450	0.8713	1.000	0.4874
Logistic_Int..	0.786	0.783	0.227	0.400	0.250	0.1540	0.8832	1.000	0.5596
Logistic_Poly	0.779	0.772	0.233	0.388	0.241	0.1503	0.8772	1.000	0.4806
RandomFor..	0.796	0.761	0.178	0.377	0.232	0.1591	0.8713	1.000	0.4115
Ridge	0.770	0.772	0.224	0.388	0.241	0.1455	0.8772	1.000	0.5778
XGBoost	0.798	0.756	0.193	0.371	0.228	0.1585	0.8683	1.000	0.0968

COMPLETE MODEL COMPARISON - 9
MODELS EVALUATED

CLASS IMBALANCE HANDLING - DETAILED IMPLEMENTATION

DATA PREPARATION

Train/Test Split (Minimal Features):

Training: 1985-01-31 to 2009-12-31 (300 samples)

Testing: 2010-01-31 to 2024-12-31 (180 samples)

Training positive rate: 22.3%


Testing positive rate: 7.2%

Imbalance handling:

- Logistic-family models: `class_weight='balanced'`
- RandomForest: `class_weight='balanced'`
- XGBoost: `scale_pos_weight = N_neg / N_pos`

MODEL BUILDING

1. Logistic Regression (Baseline)
Class-imbalance setting: `class_weight='balanced'`
AUC: 0.769
2. Ridge Logistic Regression
Class-imbalance setting: `class_weight='balanced'`
AUC: 0.770
3. Lasso Logistic Regression
Class-imbalance setting: `class_weight='balanced'`
AUC: 0.770
4. ElasticNet Logistic Regression
Class-imbalance setting: `class_weight='balanced'`
AUC: 0.770
5. Random Forest (tuned)
Class-imbalance setting: `class_weight='balanced'`
Best params: `{'bootstrap': True, 'class_weight': 'balanced', 'max_de`
AUC: 0.796
6. XGBoost
Class-imbalance setting: `scale_pos_weight=3.48`
AUC: 0.798



Model-Specific Solutions:

Logistic/Ridge/Lasso/ElasticNet: `class_weight='balanced'`

Random Forest: `class_weight='balanced'`

XGBoost: `scale_pos_weight=3.48` (exact ratio calculation)

Impact on Metrics:

Without balancing: 93% accuracy but 0% recall

With balancing: 76% accuracy but 100% recall

5. Random Forest (tuned)
Class-imbalance setting: class_weight='balanced'
Best params: {'bootstrap': True, 'class_weight': 'balanced', 'max_depth': 6, 'max_features': None, 'min_samples_leaf': 15, 'min_samples_split': 44, 'n_estimators': 271}
AUC: 0.796

Optimized Parameters

Parameter	Final Value	Search Range	Purpose
Number of Trees	271	200-1200	Model ensemble size
Max Depth	6	3-16	Prevents overfitting
Min Samples Leaf	15	1-30	Ensures leaf stability
Min Samples Split	44	2-50	Prevents over-splitting
Max Features	None (all)	[sqrt, log2, None]	Feature sampling
Bootstrap	True	[True, False]	Enables OOB scoring
Class Weight	Balanced	-	Handles imbalance

Optimization Summary

Metric	Value
Parameter Combinations Tested	60
Cross-Validation Method	TimeSeriesSplit (5 folds)
Optimization Metric	ROC-AUC
Final Test AUC	0.796
Model Rank	2nd of 9 models
Best Performer	XGBoost (0.798)


Key Insight:

The model favors conservative splitting (min_samples_split=44) and leaf size (min_samples_leaf=15), suggesting it's protecting against overfitting on the sparse recession data.

Note:

Due to sparse recession data in early validation folds, hyperparameter optimization may not have reached full potential. Production deployment would benefit from alternative validation strategies.

Random Forest Feature Importances



CUSTOM INTERACTION TERMS

Custom features created:

Original: 5

With interactions: 11

New features: 6

9. Logistic with Custom Interactions

Class-imbalance setting: `class_weight='balanced'`

AUC: 0.786

Top features by importance:


Feature	Coefficient
TS_10y3m	-0.989919
TS_10y3m_sq	-0.964517
Sahm	0.699247
Unemp_D12	0.638361
CreditSpread	0.395362
Invert_x_Unemp	0.370208
Invert_10y3m	0.277549
CreditSpread_sq	0.267863
Credit_x_Unemp	0.220649
Invert_x_Credit	0.109001

VARIANCE INFLATION FACTOR ANALYSIS


Feature	VIF	Interpretation
TS_10y3m	4.39	No issue (< 5)
CreditSpread	4.08	No issue
Sahm	3.79	No issue
Unemp_D12	2.56	No issue
Invert_10y3m	1.31	No issue

TIME SERIES CROSS-VALIDATION RESULTS

Model Performance Stability Over Time



Calibration Stability Over Time



CALIBRATION ANALYSIS - DELONG AND HOSMER-LEMESHOW TESTS

=====

FEATURE ENHANCEMENT IMPACT

=====

Logistic → Logistic_Poly: +1.3% AUC improvement
 Logistic → Logistic_Interactions: +2.2% AUC improvement
 Lasso → Lasso_Poly: +1.0% AUC improvement

=====

ADDITIONAL VALIDATION TESTS

=====

1. DELONG TEST - Statistical Comparison of Model AUCs

RandomForest (AUC=0.796) vs Logistic (AUC=0.769):
 Z-statistic: 1.360
 P-value: 0.1739
 ✗ No significant difference

RandomForest (AUC=0.796) vs XGBoost (AUC=0.798):
 Z-statistic: -0.162
 P-value: 0.8713
 ✗ No significant difference

Logistic (AUC=0.769) vs Logistic_Interactions (AUC=0.786):
 Z-statistic: -1.811
 P-value: 0.0701
 ✗ No significant difference

DeLong Test Results:

RF vs Logistic: $Z=1.36$, $p=0.174$
 (not significant)
 RF vs XGBoost: $Z=-0.16$, $p=0.871$
 (not significant)

Hosmer-Lemeshow Test:

All models $p < 0.05$ (indicating calibration issues)

Calibration Testing Performed:

Tested isotonic regression and Platt scaling methods
 Found minimal improvement (0.002 AUC gain)
 Decision: Use uncalibrated model for simplicity

Original Brier scores:

RF: 0.178 (best)
 XGBoost: 0.193
 Logistic: 0.225

2. HOSMER-LEMESHOW TEST - Calibration Quality

Logistic:
 Chi-square: 27500.102 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

Ridge:
 Chi-square: 7161.453 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

Lasso:
 Chi-square: 11067.620 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

ElasticNet:
 Chi-square: 9186.397 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

RandomForest:
 Chi-square: 244.410 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

XGBoost:
 Chi-square: 1578.376 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

Logistic_Poly:
 Chi-square: 387.679 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

Lasso_Poly:
 Chi-square: 503.529 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)


Logistic_Interactions:
 Chi-square: 3439.410 (df=8)
 P-value: 0.0000
 ✗ Poor calibration ($p < 0.05$)

PROBABILITY CALIBRATION


Calibration Analysis

- Tested isotonic regression and Platt scaling
- Found minimal improvement (0.002 AUC gain)
- Decision: Use uncalibrated model for simplicity
- Trade-off: Slightly less accurate probabilities for reduced complexity

Uncalibrated Random Forest



Calibrated Random Forest (Proper CV)



POLYNOMIAL AND INTERACTION EFFECTS

```
=====
MODELS WITH POLYNOMIAL FEATURES
=====
```

```
7. Logistic Regression with Polynomial Features
   Class-imbalance setting: class_weight='balanced'
   AUC: 0.779
```

```
8. Lasso with Polynomial Features
   Class-imbalance setting: class_weight='balanced'
   AUC: 0.778
```

```
Features selected by Lasso: 7 out of 20
```

```
Top selected features:
```

Feature	Coefficient
TS_10y3m	-1.163894
CreditSpread Sahm	0.836775
Unemp_D12	0.620179
TS_10y3m^2	-0.259923
Unemp_D12^2	-0.153324
CreditSpread	0.074585
TS_10y3m Invert_10y3m	-0.020997

3. LEAVE-ONE-RECESSION-OUT CROSS-

Logistic:

LORO-CV AUC: 0.629 ± 0.162

Number of folds: 4


RandomForest:

LORO-CV AUC: 0.613 ± 0.100

Number of folds: 4

LEAVE-ONE-
RECESSION-
OUT CROSS-
VALIDATION

CURRENT ECONOMIC ASSESSMENT - DECEMBER 2024



Current Indicator Values:

10Y-3M Spread: +0.21% (positive, not inverted)

10Y-2Y Spread: +0.33% (positive)

Credit Spread: 1.5% (moderate)

Unemployment Change: 0.0% (stable)

Model Probabilities:

Random Forest: 25.6%

XGBoost: 24.9%

Logistic: 23.2%

Historical Context:




20.57% of trading days inverted in 2024

SUMMARY ANALYSIS




	AUC	Accuracy	Precision	Recall	F1	BalancedAcc	Specificity	Brier
Logistic	0.7692307692307693	0.7611111111111111	0.23214285714285715	1.0	0.37681159420289856	0.8712574850299402	0.7425149700598802	0.22531784239235253
Ridge	0.7696913864578535	0.7722222222222223	0.24074074074074073	1.0	0.3880597014925373	0.8772455089820359	0.7544910179640718	0.2242599887943091
Lasso	0.7701520036849377	0.7722222222222223	0.24074074074074073	1.0	0.3880597014925373	0.8772455089820359	0.7544910179640718	0.22517385355912622
ElasticNet	0.7701520036849377	0.7722222222222223	0.24074074074074073	1.0	0.3880597014925373	0.8772455089820359	0.7544910179640718	0.22441084637368697
RandomForest	0.7959465684016582	0.7611111111111111	0.23214285714285715	1.0	0.37681159420289856	0.8712574850299402	0.7425149700598802	0.17847806571518396
XGBoost	0.7977890373099953	0.7555555555555555	0.22807017543859648	1.0	0.37142857142857144	0.8682634730538923	0.7365269461077845	0.192864725220321
Logistic_Poly	0.7789037309995394	0.7722222222222223	0.24074074074074073	1.0	0.3880597014925373	0.8772455089820359	0.7544910179640718	0.23316955847885426
Lasso_Poly	0.7779824965453708	0.7722222222222223	0.24074074074074073	1.0	0.3880597014925373	0.8772455089820359	0.7544910179640718	0.2257800056924869
Logistic_Interactions	0.7858129894058038	0.7833333333333333	0.25	1.0	0.4	0.8832335329341318	0.7664670658682635	0.2272009476308868

Key Performance Metrics



CURRENT RISK ASSESSMENT & 24-MONTH FORECAST

Forecast: Model projects moderate recession risk (~25–30% probability) through late 2025, rising slightly by 2026, but no clear imminent recession signal as of Dec 2024.



BUSINESS IMPACT & APPLICATIONS



Applications:

Portfolio Management: Risk-based asset allocation

Corporate Planning: Capital expenditure timing

Banking: Loan provisioning and stress testing

Policy Making: Early intervention opportunities

KEY CONTRIBUTIONS & NOVELTY

Technical Contributions

First **multi-indicator framework**: 7 Treasury spreads (lead-lag + logistic baselines) + credit spreads + unemployment dynamics

Explicit **class imbalance correction** across all models (balanced weights, exact XGBoost ratio)

Advanced Random Forest tuning with time-series CV → lifted AUC by ~5%

Unified evaluation framework: ROC/PR curves, DeLong tests, Hosmer–Lemeshow calibration, Brier scores

Practical Impact

~82% accuracy & strong recall at 12-month horizon

Improved calibration & interpretability (SHAP, feature importance)

Holistic economic view: integrates financial market and labor indicators for more robust forecasting

Deployment-ready pipeline: fully automated, reproducible, extendable to new data

CONCLUSIONS & RECOMMENDATIONS



Confirmed Findings:

- Yield curve inversions remain highly reliable recession predictors
- The **10Y–3M spread** is the single most powerful signal
- A **multi-indicator framework** (adding credit spreads & unemployment dynamics) improves accuracy and reduces false alarms

Recommendations:

- **Monitor** closely when probability $> 30\%$ (early warning)
- **Take defensive action** when probability $> 50\%$ (elevated risk)
- Always **combine model signals with expert judgment** and real-time context

THANK you!

THANK YOU FOR YOUR ATTENTION

GitHub Repository: <https://github.com/BimboSolanke/End-to-end-ML-pipeline-for-predicting-recessions-using-yield-curves-and-macroeconomic-analysis>